

Child Care Subsidies with One- and Two-Parent Families

Emily G. Moschini *

June 27, 2020

Abstract

This paper examines the role of family structure in determining the effect of child care subsidies on child skill accumulation. Family structures differ in the number of parents and in the way parental time and child care time combine to produce investment in child skill. Using a nationally representative panel data set which follows families raising young children in the United States, I estimate skill investment functions and establish that one- and two-parent families use different technologies to invest in their children. In particular, I find that time in child care has a larger effect on child skill for children of one-parent families. To analyze the implications for policy, I nest the family investment problems in a general equilibrium model in which one- and two-parent families form endogenously. I calibrate this model to match empirical moments of the relationship between family income, family structure, and child skill, and then use it to examine the equilibrium effects of a universal child care subsidy. I find that a subsidy improves skill outcomes for all children, thus increases aggregate welfare by up to 5.6%, and reduces inequality in child outcomes by having a larger effect on children of one-parent, low-income families. To separate the roles of technology and parent composition in generating the gap in child skill outcomes at any subsidy level, I construct a counterfactual in which one- and two-parent families form according to a random shock. I find that between 28% and 40% of the gap in child skill outcomes across family structures is driven by technological differences, depending on the level of the child care subsidy.

Keywords: Early Childhood Education, Family Structure, Households, Skills

JEL Classification Numbers: J24, J12, I21, D13

*William & Mary. *email:* egmoschini@wm.edu. *Acknowledgements:* I would like to thank Anmol Bhandari, Adam Blandin, Elizabeth Caucutt, Mariacristina De Nardi, Diego Daruich, Eric French, Christopher Herrington, Loukas Karabarbounis, Karen Kopecky, Jeremy Lise, Lance Lochner, Ellen McGrattan, Peter McHenry, Joseph Mullins, and Sergio Ocampo Díaz, as well as seminar participants at Virginia Commonwealth University, Midwest Macro (Athens, GA), the SED (St. Louis, MO), and the 2020 Barcelona Summer Forum for comments. All errors are mine.

1 Introduction

There is widespread consensus that early childhood is a key phase in the development of skill, and that policies targeted at this age group can improve educational attainment and labor market outcomes later in life (see Almond and Currie (2011) for a review).¹ This consensus is founded on evaluations of several small-scale randomized controlled trials (RCTs) with long-run follow-ups, where children under the age of 5 from impoverished backgrounds are exposed to active engagement with adults in environments structured to foster the development of child skill. Participation in these programs has been linked with long-run improvements in educational achievement and adult earnings (see Heckman, Stixrud, and Urzua (2006), Anderson (2008), Cunha, Heckman, and Schennach (2010), and Garcia, Heckman, and Ziff (2017), among many).

The evidence from RCTs is complemented by studies of large-scale subsidized child care programs with universal eligibility, for example Baker, Gruber, and Milligan (2008), Berlinski, Galiani, and Gertler (2009), and Gupta and Simonsen (2010).² Recent evaluations of such programs have emphasized heterogeneity in the effects of child care subsidies on child outcomes across demographic attributes. In particular, Havnes and Mogstad (2014) use Norwegian data to show that the effects on lifetime earnings are larger for children of poor families than those of rich families, and Kottelenberg and Lehrer (2017) use data from Canada to show that the effects on child development scores are higher for children of single mothers than for children of couples. In the United States, these demographic attributes are quantitatively relevant and intertwined: single mothers are much more likely to be poor than parenting couples³, and they raise 20 percent of children under 5.

In this paper I examine the role of family structure in determining the effect of child care subsidies on child skill accumulation. The way in which families allocate their children's time differs systematically across family structures: I find that one-parent families use more child care, and contribute less quality time from parents, than couples do (despite facing similar relative prices). These choices determine the activities and environments that children experience, and I interpret them as reflecting an investment decision made by families.

Building on the observed disparities in these family choices, I first establish that there is a significant difference in how one- and two-parent families combine their own time and purchased child care to generate investment in their children. Next, I incorporate investment in child skill into a general equilibrium framework with endogenous family formation and evaluate the effects of a universal child care subsidy in this setting. Without a subsidy, children of one-parent families have lower levels of skill at the beginning of adulthood than children raised by couples. I find that the subsidy increases child skill outcomes, especially for one-parent families, thereby narrowing the gap across family structures. The model also allows me to

¹For the perspective of child psychologists, see Shonkoff (2010) and Nisbett, Aronson, Blair, Dickens, Flynn, Halpern, and Turkheimer (2012).

²Examples include programs in Oklahoma (started in 1998), Quebec (1997), Argentina (1993), Norway (1975), and Denmark (1964).

³In the 2000 U.S. census, single mothers with children under 5 had a poverty rate of 41 percent, versus 6 percent for married couples (author's calculations).

decompose the gap into a composition portion and a portion driven by investment technologies. I find that heterogeneity in investment technologies drives a sizeable portion of the gap.

To establish that one- and two-parent families use different investment technologies, I analyze the problem of families investing in their children's skill by choosing time from parents and child care bought on the market. I allow the parameters which govern how investment is generated to differ across family structures. I then use the problem of the parents to derive estimation equations for these parameters, and test whether the two technologies are different. I find that they are.

I construct the implied price of investment for each family structure, and derive the elasticity of the price of investment with respect to the price of child care, in order to highlight why heterogeneity in investment technologies matters for policy design. A child care subsidy will increase children's skill outcomes to the extent that it decreases the price of investment in skill. Consistent with research that documents higher gains for children of one-parent families, I find in the data that this elasticity is three times higher for one-parent families than it is for two-parent families. I believe this provides a novel perspective on how investment technologies mediate the response of parental investments to child care subsidies.

My main data source is the Early Childhood Longitudinal Study, Birth Cohort (ECLS-B), a family-level data set from the US Department of Education (Department of Education (2001)). Unlike other data sets commonly used to estimate skill accumulation technologies during early childhood, the ECLS-B is designed to be representative of families raising 9-month old children in the United States.⁴ The ECLS-B provides information on parental and market time inputs, parental wages, child care prices, and family structure for children over three waves of the survey, corresponding to when the child is 9 months, 2 years, and 4 years old. This allows me to control for unobserved parenting productivities by implementing a fixed effects estimator. In addition, I observe measures of child skill in each of these waves. Together with family attributes, these test scores generate empirical moments which I use to calibrate the general equilibrium framework.

Evaluating the implications of heterogeneity in skill investment technologies for the effects of child care policies requires a larger model, where both the marginal cost but the marginal benefit of skill adjust to the subsidy in equilibrium. The marginal benefit of skill operates through better marriage prospects and increased earnings. As a child care subsidy makes skill investment cheaper, families may form differently when the costs of parenting change (thus making skill more or less useful on the marriage market), or the labor market return to skill may decrease as labor income taxes increase to finance the subsidy (which would dampen increases in investment due to the subsidy).

To incorporate these forces, I analyze the effects of child care subsidies in the general equilibrium of an expanded model. This environment features a marriage market, in which each generation of men and women meets via random search and endogenously forms one- or two-parent families. After forming, families spend a portion of their lifetime altruistically investing in their children's skill using the technology of their family structure. This phase of life matches the optimization problem used in the estimation. Skill investment is the

⁴Two common sources of data used in similar studies are the Panel Study of Income Dynamics, Child Development Supplement (PSID CDS) and the Children of the National Longitudinal Survey of Youth 1979 cohort (NLSY79 Child).

only form of transfer across generations—at the end of life, families die without leaving financial bequests.

I calibrate the model to match empirical moments of child skill, family income, and family formation. The model also qualitatively matches untargeted empirical facts on the gap in child outcomes across family structures, the positive assortative matching by skill within couples, and the marriage rate over the family income distribution. In addition, the model output quantitatively aligns with empirical measures of the elasticity of the price of investment with respect to the subsidy, computed from the ECLS-B using the estimated investment technologies.

I examine the equilibrium effects of a universal child care subsidy with the fully parameterized model. I find that such a subsidy increases aggregate welfare by up to 5.6% in consumption equivalent units. It does so by mimicking the contract which children would like to make with their parents and others in society, promising a portion of their future earnings to fund investment in their skill during childhood. The subsidy makes investment cheaper, and then taxes future earnings to recoup the cost.

I find that my model framework delivers a gap in child skill outcomes between children of one- and two-parent families, which is reduced by the child care subsidy. This is qualitatively in line with the findings of Kottelenberg and Lehrer (2017). Although the child care subsidy is the same level for all families, the effect it has on parental investment in children is heterogeneous: the subsidy has a larger effect on children of one-parent, low-income families. Both endogenous and exogenous differences across family structures cause this gap. Besides differences in income composition due to endogenous family formation, the two family structures differ in how their time inputs affect their children's skill.

Using the model, I carry out a counterfactual exercise in which I seek to separate the role of family income from family structure in determining child skill outcomes. I then examine how this decomposition changes in response to a child care subsidy. To do this, I construct a child outcomes in an environment in which one- and two-parent families form exogenously, by assign family structure randomly while keeping the marriage rate the same. The resulting counterfactual gap in child skill outcomes is driven entirely by differences in exogenous attributes of family structure (in particular, investment technologies), while the difference between the equilibrium gap and counterfactual gap is driven by the relative composition of parental skill and family income in one-parent families. I find that between 28% and 40% of the gap in child skill outcomes across family structures is driven by exogenous factors, depending on the level of the child care subsidy.

This study features altruistic parents investing in the skill of their young children, a framework for thinking about intergenerational transfers which dates back to Barro and Becker (1989). I have this in common with both Lee and Seshadri (2019) and Daruich (2020), the two studies most closely related to mine. However, these papers examine different aspects of subsidies to skill accumulation. Lee and Seshadri (2019), for example, compares skill subsidies at different stages of life.⁵ Daruich (2020) examines the effects of transfers-in-kind child care subsidies on the aggregate economy. The focus of that study is on the scaling up of free high-quality child care programs, and the economy's resulting transition to a new equilibrium.

⁵ Another paper with this approach is Caucutt and Lochner (2020), which emphasizes that the timing of binding borrowing constraints determines the optimal timing for interventions. Restuccia and Urrutia (2004) instead emphasize the relative importance of early parental investments in determining adult outcomes and the implications of this for broader education policy design.

Broadly, these analyses conclude that families play a large role in shaping adult outcomes and that the earlier in life the subsidy is received, the better. These results are consistent with findings from Cunha, Heckman, and Schennach (2010), and Agostinelli and Wiswall (2017). Building on these results, I narrow skill investment decisions to being made only by parents, and consider only subsidies during early childhood.

Another important feature of my model is endogenous family formation, where both one- and two-parent families exist in equilibrium. Other macroeconomic studies also incorporate endogenous family formation into their frameworks, for example Abbott, Gallipoli, Meghir, and Violante (2019), which focuses on college financial aid in the United States⁶, and Gayle, Limor, and Soytas (2017), which seeks to attribute the observed intergenerational correlation earnings to various sources (including parental investments). However, the frameworks of these papers do not feature multiple family structures in equilibrium. Instead, the authors emphasize positive assortative matching of couples as a magnifying force for inequality.

In my model I do not focus on analyzing the labor force participation of parents. The extensive margin could be crucial for my policy analysis if the purpose of child care subsidies was to affect the labor supply and earnings of parents. A study which focuses on this potential effect is Guner, Kaygusuz, and Ventura (2016), who allow for heterogeneity in family structures and consider how this heterogeneity leads to different policy responses. However, the focus is on the labor supply of parents, not the skill accumulation of children, and family structure is specified to emphasize that decision. My specification explicitly emphasizes investment in child skill rather than how having a spouse affects labor supply choices of adults.

In my framework, the critical primitive is the function that maps from family time use decisions to child skill outcomes. The distinctive feature of my specification for this technology is that I only allow for time inputs (parental and non-parental time), and the fact that it incorporates a possibly binding time constraint on the child. Concerning the inclusion of both parents for the couple's technology, there are several papers that estimate an investment technology for couples. These include Abbott (2020) and Del Boca, Flinn, and Wiswall (2014), although the investment inputs in those specifications differ from mine. Those studies also differ in their emphasis on income risk (Abbott (2020)) and the changing role of mothers at each stage of early childhood (Del Boca, Flinn, and Wiswall (2014)), which are reflected in their specifications.

My derivation of estimation equations is similar to that of Lee and Seshadri (2019), despite different specifications for the skill accumulation process. In particular, I do not include multiple phases of investment, each with different technology parameters. Instead, I model multiple family structures with different parameters, and only one phase of investment—early childhood. Daruich (2020) calibrates an investment technology within his model. In his specification, parent time and money (goods) are combined to form investment, and subsidized child care is like a transfer-in-kind of the goods input. Unlike his specification, I use expenditures on child care time as an alternative input to parental time, which reflects the fact that in the data money spent on young children is mostly spent on child care.⁷

The paper proceeds as follows. Section 2 presents the investment problem of one- and two-parent families,

⁶Lochner and Monge-Naranjo (2011) also focus on this policy, although families do not form endogenously in that study.

⁷My calculations supporting this claim use the PSID CDS and are reported in the appendix.

which form the core of my framework. From these problems I derive estimation equations. In Section 3, I present the data used to estimate the parameters of the skill investment technologies, and the estimation results. Section 4 discusses how heterogeneity in investment technologies matters for investment's response to a child care subsidy. In this section I derive elasticities, both for investment's price and (given a utility function) for the marginal cost of investment, with respect to the price of child care. Section 5 presents the general equilibrium model framework, model parameterization, and model fit. Section 6 reports the child care policy experiment and decomposition results. Section 7 concludes.

2 Investment in Child Skill

The critical family decision that identifies parameters of the skill investment technologies is the allocation of child time between child care and time with parents. I interpret this decision as determining investment in child skill, and thereby child skill accumulation. When parents choose how to allocate their child's time, they internalize how these activities and environments affect child skill, while balancing the relative prices and productivities of these different sources of inputs. For example, parents make the decision of how much time to personally spend with their child, doing activities like reading to them and playing with them outside, by weighing the necessary foregone earnings against the cost of using child care instead. If the child care is very high quality, this lowers the cost of using it as a substitute for parental time, and vice versa.

In this section, I formalize the problems of parents in one- and two-parent families. The purpose of this section is to highlight how parents combine their own time and child care time to produce investment in their child's skill, and how this decision depends on attributes of their investment technology. The optimization conditions for the parenting problems are used to derive estimation equations for the investment technology parameters.

In what follows, I refer to one-parent families as single mothers and two-parent families as couples. This aligns with the observed characteristics of one- and two-parent families: one-parent families are by far mostly single mothers, and two-parent families may be either married or cohabiting couples.

2.1 Parenting Problems

Parents make skill investment decisions for their child for each of the J periods of childhood. The child is born with some initial skill θ_1 . In each period t , the parent affects the child's stock of skill in that period, θ_t , by their choice of investment, I_t^{type} , according to the skill production function $\theta_{t+1} = f(\theta_t, I_t^{type})$, where $type \in \{SM, MC\}$ for single mothers (SM) and married couples (MC), respectively.

Investment is generated from parental time and time purchased on the market in the form of non-parental child care. Time contributed from parent type j is denoted $q_{j,t}$, where j can refer to a single mother ($j = 1$), a married or cohabiting mother ($j = 2$), or a married or cohabiting father ($j = 3$). Child care time is denoted n_t . The price of parental time is the wage of that parent, and the price of child care time is denoted p_t .

Single mothers are indexed by the initial skill of their child, θ_1 , the productivity of child care they use, ϕ_n ,

the parenting productivity ϕ_1 , and their wages in each time period, $\{w_{1,t}\}$. During parenthood, a single mother chooses non-parental child care time n_t , and her own time investments in her child $q_{1,t}$, in each period, to produce her child's skill accumulation. The single mother also chooses consumption, c_t , and savings, a_{t+1} , to solve the following problem:

$$\begin{aligned} \max_{\{c_t, \ell_t, a_{t+1}, n_t, q_{1,t}\}_{t=1}^J} & \quad \left[\sum_{t=1}^J \beta^{t-1} u^{SM}(c_t, \ell_t) \right] + \beta^{J-1} bV^{child}(\theta_{J+1}) \\ \text{s.t. } \theta_{t+1} &= f(\theta_t, I_t^{SM}(\phi_n n_t, \phi_1 q_{1,t})) \\ c_t + a_{t+1} + p_t n_t &\leq w_{1,t}(1 - \ell_t - q_{1,t}) + (1 + r_t) a_t + T, \quad \forall t \end{aligned} \quad (1)$$

where, in each period, $n_t, \ell_t, q_{1,t} \geq 0$, and $a_1, a_{J+1} = 0$. The time constraints for the mother and child are $\ell_t + q_{1,t} \leq 1$ and $n_t + q_{1,t} \leq 1$, respectively. Here, T is the lump sum transfer the single mother receives from the government, and $1 + r_t$ is the return on savings. This formulation assumes that investment in children is motivated by altruistic returns to the parent from the child's expected lifetime utility, expressed as a function of their skill at adulthood.⁸ Parents can affect this outcome through the skill accumulation technology. Specifically, parental and non-parental child care time generate investment in child skill according to $I_t^{SM}(\phi_n n_t, \phi_1 q_{1,t})$.

Married couples are similarly indexed by the initial skill of their child, θ_1 , the productivity of child care ϕ_n , the parenting productivity of the female and male parent, ϕ_2 and ϕ_3 , and the wages of each parent in each period, $\{w_{2,t}, w_{3,t}\}_{t=1}^J$. Both parents and the child have a one-unit time endowment in each period. The couple jointly chooses consumption, c_t , savings, a_{t+1} , child care purchased on the market, n_t , and quality time from the female and male parents, $q_{2,t}$ and $q_{3,t}$, to solve the following problem:

$$\begin{aligned} \max_{\{c_t, \ell_t, a_{t+1}, n_t, q_{2,t}, q_{3,t}\}_{t=1}^J} & \quad \left[\sum_{t=1}^J \beta^{t-1} u^{MC}(c_t, \ell_t) \right] + \beta^{J-1} bV^{child}(\theta_{J+1}) \\ \text{s.t. } \theta_{t+1} &= f(\theta_t, I_t^{MC}(\phi_n n_t, \phi_2 q_{2,t}, \phi_3 q_{3,t})) \\ c_t + a_{t+1} + p_t n_t &\leq w_{2,t}(1 - \ell_t - q_{2,t}) + w_{3,t}(1 - \ell_t - q_{3,t}) + (1 + r_t) a_t + T, \quad \forall t \end{aligned} \quad (2)$$

where, in each period, $n_t, \ell_t, q_{2,t}, q_{3,t} \geq 0$, $a_1, a_{J+1} = 0$, and the time constraints are $\ell_t + q_{2,t} \leq 1$, $\ell_t + q_{3,t} \leq 1$, and $n_t + \max\{q_{2,t}, q_{3,t}\} \leq 1$. The last constraint embeds the assumption that parents are allowed to invest in the child at the same time (their quality time investments are non-rival).

2.2 Cost Minimization

Implicit in these optimization problems are intra-temporal cost-minimizing problems, where families select the cheapest combination of inputs to finance a desired level of investment, subject to their skill investment technology.

⁸For a discussion of this modelling choice, see Appendix II.

For single mothers, for a given level of investment I_t the time inputs from the mother and child care provider solve:

$$\begin{aligned} \min_{\{q_{1,t}, n_t\}} \quad & w_{1,t}q_{1,t} + p_t n_t \\ \text{s.t.} \quad & I_t^{SM}(\phi_n n_t, \phi_1 q_{1,t}) \geq I_t \end{aligned} \quad (3)$$

where the time constraint $n_t + q_{1,t} \leq 1$ is assumed not binding.

In what follows I will specify a linear homogeneous investment function $I_t^{SM}(\cdot)$, implying that the value function of the cost minimization problem satisfies $C(w_{1,t}, p_t; I_t) = I_t \Lambda(w_{1,t}, p_t)$. Hence, $\Lambda(w_{1,t}, p_t)$ can be interpreted as the implicit price index of investment in period t (a composite of the prices of the two inputs, the mother's wage and the price of child care).

Similarly, for a given level of investment I_t couples choose inputs $\{q_{2,t}, q_{3,t}, n_t\}$ to solve:

$$\begin{aligned} \min_{\{q_{2,t}, q_{3,t}, n_t\}} \quad & w_{2,t}q_{2,t} + w_{3,t}q_{3,t} + p_t n_t \\ \text{s.t.} \quad & I_t^{MC}(\phi_n n_t, \phi_2 q_{2,t}, \phi_3 q_{3,t}) \geq I_t \end{aligned} \quad (4)$$

where the time constraint $n_t + \max\{q_{2,t}, q_{3,t}\} \leq 1$ is assumed to be not binding. Again, provided the investment function is homogeneous of degree one, the value function of the cost minimization problem satisfies $C(w_{2,t}, w_{3,t}, p_t; I_t) = I_t \Lambda(w_{2,t}, w_{3,t}, p_t)$, where $\Lambda(w_{2,t}, w_{3,t}, p_t)$ is the relevant price of investment for the married couple (a composite of the prices of the three inputs).

2.3 Optimal Time Choices

The skill investment functions for the two family types are assumed to have the Constant Elasticity of Substitution (CES) structure. Specifically, for single mothers and married couples, they are, respectively:

$$I_t^{SM} = \left[\alpha_1 (\phi_1 q_{1,t})^{\frac{\eta_1-1}{\eta_1}} + (1 - \alpha_1) (\phi_n n_t)^{\frac{\eta_1-1}{\eta_1}} \right]^{\left(\frac{\eta_1}{\eta_1-1} \right)} \quad (5)$$

$$I_t^{MC} = \left[\alpha_2 (\phi_2 q_{2,t})^{\frac{\eta_2-1}{\eta_2}} + \alpha_3 (\phi_3 q_{3,t})^{\frac{\eta_2-1}{\eta_2}} + (1 - \alpha_2 - \alpha_3) (\phi_n n_t)^{\frac{\eta_2-1}{\eta_2}} \right]^{\left(\frac{\eta_2}{\eta_2-1} \right)} \quad (6)$$

Here, η_1 is the elasticity of substitution for the single-mother function, η_2 is the elasticity of substitution for the married couple function, and α_j ($j = 1, 2, 3$) are share parameters. The heterogeneity in the technology of skill investment functions is apparent in the structure of these two functions: for single mothers, only two inputs are available, where married couple can draw on three distinct inputs. Note, however, that the SM technology is nested in the MC technology, a feature of this specification that will be used for hypothesis testing below. Specifically, the MC skill investment function reduces to the SM function if $\alpha_2 = \alpha_1$,

$\alpha_3 = 0$, and $\eta_2 = \eta_1$.

To estimate the parameters of the skill investment functions, I rely on the optimality conditions of the cost minimization problems discussed earlier. Specifically, the cost-minimizing solutions for the *SM* problem, $q_{1,t}^*$ and n_t^* , satisfy the tangency condition:

$$\frac{\frac{\partial I_t^{SM}(\phi_1 q_{1,t}^*, \phi_n n_t^*)}{\partial n_t}}{\frac{\partial I_t^{SM}(\phi_1 q_{1,t}^*, \phi_n n_t^*)}{\partial q_{1,t}}} = \frac{p_t}{w_{1,t}}$$

Similarly, the optimal solutions of the *MC* cost minimization problem, $q_{1,t}^*$, $q_{3,t}^*$ and n_t^* , satisfy:

$$\frac{\frac{\partial I_t^{MC}(\phi_2 q_{2,t}^*, \phi_3 q_{3,t}^*, \phi_n n_t^*)}{\partial n_t}}{\frac{\partial I_t^{MC}(\phi_2 q_{2,t}^*, \phi_3 q_{3,t}^*, \phi_n n_t^*)}{\partial q_{j,t}}} = \frac{p_t}{w_{j,t}} \quad j = 2, 3$$

The optimality conditions for the two family structures can then be written as:

$$\ln \left(\frac{q_{1,t}^*}{n_t^*} \right) = \eta_1 \ln \left(\frac{\alpha_1}{1 - \alpha_1} \right) - \eta_1 \ln \left(\frac{w_{1,t}}{p_t} \right) - (\eta_1 - 1) \ln \left(\frac{\phi_n}{\phi_1} \right) \quad (7)$$

$$\ln \left(\frac{q_{j,t}^*}{n_t^*} \right) = \eta_2 \ln \left(\frac{\alpha_j}{1 - \alpha_2 - \alpha_3} \right) - \eta_2 \ln \left(\frac{w_{j,t}}{p_t} \right) + \left(\frac{1}{\eta_2 - 1} \right) \ln \left(\frac{\phi_j}{\phi_n} \right) \quad j = 2, 3 \quad (8)$$

In the next section, I describe how these equations are applied to the data to find point estimates for the parameters of the two skill investment technologies.

3 Estimation

The cost-minimizing solutions of equations (7) and (8) provide the structural representation for a panel regression with which I estimate the investment technology parameters. Specifically, I observe the relevant variables measured in each period t for each family i with parental type j (recall that $j = 1$ for a single mother, $j = 2$ for a married or cohabiting mother, and $j = 3$ for a married or cohabiting father). Hence, the estimation equations are written as:

$$y_{j,i,t} = \beta_{0,j} + \beta_j x_{j,i,t} + \tau_{j,i} + \nu_{j,i,t} \quad (9)$$

Here, the dependent variable is $y_{j,i,t} \equiv \ln \left(\frac{q_{j,i,t}}{n_{j,i,t}} \right)$, the independent variable is $x_{j,i,t} \equiv \ln \left(\frac{w_{j,i,t}}{p_{j,i,t}^n} \right)$, and the residual error term $\nu_{j,i,t}$ is assumed to be i.i.d. Note that, in view of (8), the slope parameters for the *MC* parents are the same, that is $\beta_2 = \beta_3$ —a restriction that is maintained in estimation.

In the foregoing panel regression, $\tau_{j,i}$ captures the unobserved and time-invariant relative efficiency coefficients of inputs, which are evident in equations (7) and (8). If the investment productivities of parental time

and child care time are correlated with the observed prices of these inputs, which seems plausible, then the coefficient on the price ratio regressor will be biased in an Ordinary Least Squares (OLS) estimation because of the omitted variable in the residual. In light of this, I estimate the model using a Fixed Effects estimation at the parent level for these equations. That is, I treat $\tau_{j,i}$ as fixed effects to be estimated and, by a suitable normalization, I will assume $E[\tau_{1,i}] = E[\tau_{2,i}] = E[\tau_{3,i}] = 0$. This provides the family-specific intercepts $\beta_{0,j}$ that are needed to retrieve the structural CES parameters. Specifically, the parameters of the estimating equation map into the structural parameters of interest as follows:

$$\eta_j = -\beta_j, \quad j = 1, 2 \quad (10)$$

$$\alpha_1 = \frac{\exp\left(-\frac{\beta_{0,1}}{\beta_1}\right)}{1 + \exp\left(-\frac{\beta_{0,1}}{\beta_1}\right)} \quad (11)$$

$$\alpha_j = \frac{\exp\left(-\frac{\beta_{0,j}}{\beta_2}\right)}{1 + \exp\left(-\frac{\beta_{0,2}}{\beta_2}\right) + \exp\left(-\frac{\beta_{0,3}}{\beta_2}\right)}, \quad j = 2, 3 \quad (12)$$

3.1 Data

Two datasets are combined to measure parental educational time inputs ($q_{j,i,t}$), non-parental child care time inputs ($n_{j,i,t}$), hourly wages ($w_{j,i,t}$), and hourly non-parental child care prices ($p_{j,i,t}$) (recall that the indexes denote parental type j in family i in period t). I use the ECLS-B, which is a panel data set, and the American Time Use Survey (ATUS), which is a repeated cross-section sampled from the Current Population Survey (CPS). Data for hourly wages, hourly price of child care, quality time from the parents in weekly frequency of activities, and non-parental child care time in hours per week can be constructed from information reported in the ECLS-B. The ATUS contributes measures of hours per activity, which I impute to the ECLS-B to transform frequencies into hours per week of parental educational time inputs.

The Early Childhood Longitudinal Study, Birth Cohort The ECLS-B reports labor earnings, the period of time over which the labor earnings were accrued (a day, a week, two weeks, etc.), and the hours worked in a week separately for each parent.⁹

After-tax labor earnings are converted into hourly wages using hours worked per week.¹⁰ If hours worked were not reported for the parent, they are imputed using the response to part-time or full-time status (assigning 30 or 40 hours worked per week, respectively). The result is hourly after-tax wages for mothers and fathers, conditional on observing labor earnings and some information about the intensity of labor supply.¹¹

⁹In addition, there are assessments of child skill reported in each wave of the survey. These are used when I compare the model's moments to untargeted moments in the data. For more information on child skill measures in the ECLS-B, see Appendix I.

¹⁰To account for taxes, labor earnings are corrected using the slopes (tax rates) from Table 2 of McGrattan and Prescott (2017). This adjustment accounts for the progressive nature of the US tax system.

¹¹I do not impute hourly wages for observations without this information.

To calculate hourly price of non-parental child care, I use spending on the primary source of non-parental child care, adjusted by the number of weeks that cost represents and the hours per week the source of child care is used by the family.¹²

The resulting price per hour for non-parental child care, and hourly wages of mothers and fathers, are observed for each family in the first three waves of the survey when the variables necessary to construct them are reported.¹³

Hours purchased for the primary source of child care are reported directly in the ECLS-B. Quantities of parental educational time, the other source of investment inputs, are constructed using information reported in both the ECLS-B and the ATUS. Various activities that parents do with their children are reported in the first three waves of the ECLS-B; these activities are reported in units of frequencies (every day, once a week, etc.). Only some of these activities are considered to be educational time: these are activities with the child that include talking and reading and spending time outside.¹⁴ This definition is consistent with the literature for the importance of active time with children (Del Boca, Flinn, and Wiswall (2014)). In order to convert observed quality time from frequencies into hours per week, the next step is to impute time per activity from the ATUS.

The American Time Use Survey Data on levels of time per activity for a parent with a given set of characteristics come from the 2003-2016 pooled ATUS sample. This dataset provides a time diary along with CPS variables on age, gender, marital status, labor force status, educational attainment, parental status, and child age. Individuals are restricted to be between 15 and 55 years of age, with a child 3 years or younger. I use information on gender, marital status (married/cohabiting or single), labor force status (participating or not), and educational attainment, where educational attainment is discretized into those with a high school degree or less, and those with more than a high school degree. I calculate the survey-weighted average of time spent on an activity (conditional on engaging in it) for each group, for both time spent reading to the child and time spent playing with the child.¹⁵

Imputation After linking parents in the ECLS-B with their appropriate group in the ATUS along dimensions of age, gender, marital status, and educational status, the ATUS levels of time spent reading and time spent playing are assigned to reading activities and playing outside activities in the ECLS-B. All ECLS-B observations with the same age, gender, marital status, and education bin are assigned the same number of hours to each time they report engaging in an activity with their child. Next, total quality time per parent in each family in each wave is calculated by summing across the two activities.

¹²The ECLS-B also reports all sources of child care; the relevant moments for the calibration are not affected by using total hours in child care instead of only the primary source.

¹³To see how the sample selection procedure affected the composition of child care sources used by families, see part I of the Appendix.

¹⁴This definition is founded in part on an exercise in which the definition of quality time was varied to identify which definition led the model estimation equations to best fit the data.

¹⁵Summary statistics and tabulations of the ATUS sample, along with summary statistics of the raw ECLS-B sample, are reported in Appendix I.

Restrictions on the Estimation Sample The estimation sample for mothers and fathers parenting in couples pools those who are married and those who are cohabiting. The single mother estimation sample is composed of mothers in the sample who are a primary caregiver and who do not have a significant other living in the household with them. For both couples and single mothers, observations are only admissible if the resident primary caregiver is a biological parent, is less than 55 years of age, reports working for pay, makes less than 200 dollars an hour, pays at least 1 cent per hour for the primary source of child care, and whose child spends at least 0.1 hours per week in child care. Families are only valid observations if the biological mother had her first child after age 15 and before age 45. Finally, I only use families I observe in all of the first three waves of the ECLS-B. This leaves me with 50 single mothers and 300 couples (with sample sizes rounded to the nearest 50, as per NCES requirements).

In the sample selection I restrict attention to families for which the estimation equation is valid. With regards to non-participation, I seek to use families that satisfy the assumptions of my estimation equation. These families I interpret as obeying the tangency conditions of the cost-minimization problem. If a parent did not participate in the workforce, or if the child's time constraint was binding, the estimation equation would not hold, although I assume the investment technology would be the same.

Estimation Sample Moments Moments from the estimation samples for each parental type are presented in Table 1. By comparing across parental types, the following three main qualitative points are apparent, related to quantities and relative prices of investment inputs as well as family income.

First, couples contribute more parental time, and purchase less child care, than single mothers do. Mothers in a couple spend on average 9 hours per week engaging in educational activities with their children, while fathers spend an average of about 5 hours per week, and an average couple purchases about 33 hours per week of child care time. From Table 1, single mothers spend on average almost 2.5 hours in educational activities with their children, and purchase 36 hours per week of non-parental child care time. The ratio of mother's time to non-parental child care time is correspondingly lower than for married mothers, at 0.09 for the former compared to 0.37 for the latter.

Second, single mothers make on average about seven dollars per hour less than mothers parenting in couples. However, single mothers tend to use cheaper child care than couples do, so that the ratio of the price of child care to the mothers hourly wage is about the same as it is for married and cohabiting mothers, at 35% relative to 39% for the latter.

Finally, single mothers have an average income less than half that of married mothers or fathers. Differences in income are echoed by differences in other socioeconomic attributes across the two family structures. Starting with couples, both parents are on average in their mid-30s, with mothers who are married or cohabiting having had their first child at the age of 29 on average. Parents raising young children in a couple are well-educated: 65% of mothers have a bachelors degree or higher, and 54% of fathers do. Single mothers, by comparison, are less educated than parents raising children in a couple (21% of single mothers have a college degree) and are younger both at the time of data collection and at the age at which they had their first child. Men and women parenting in couples are not poor: only 1% of mothers in the sample are below

the poverty line (0% of the fathers are), while only 6% of mothers and 5% of fathers are below 185% of that threshold.¹⁶ The poverty rate of single mothers is nearly thirty time higher than couples, at 28%, while the percent of single mothers below 185% of the poverty line is more than eight times higher, at 51%.¹⁷

Table 1: Estimation Sample Moments (Waves 1-3, Unweighted)

	Singles: Mothers			Couples: Mothers			Couples: Fathers		
	mean	p50	sd	mean	p50	sd	mean	p50	sd
Parental Time	2.53	2.74	1.13	9.33	9.05	6.84	4.80	4.72	2.39
Child Care Time	36.61	40.00	12.10	33.96	40.00	12.54	33.38	37.00	12.58
Parental Hourly Wage	8.72	7.91	4.16	15.80	12.70	13.00	16.11	12.77	11.19
CC Hourly Price	2.77	2.00	3.08	4.09	3.05	4.24	4.06	3.00	4.34
Ratio (Prices)	0.35	0.22	0.43	0.39	0.24	2.80	0.31	0.22	0.57
Ratio (Quantities)	0.09	0.07	0.11	0.37	0.23	0.56	0.20	0.14	0.22
Family Income	35989	27500	40162	106711	87500	67092	106425	87500	64504
Family Income (After Tax)	39527	34021	28055	88719	75730	46134	88549	75730	44307
Age of Parent	28.76	28.00	5.82	34.17	34.00	5.16	35.63	35.00	5.45
Age First Child (Mother)	23.61	22.50	5.42	29.04	29.00	5.37			
BA or higher	0.21			0.65			0.54		
Poor (100% Pov. Line)	0.28			0.01			0.00		
Poor (185% Pov. Line)	0.51			0.06			0.05		
Observations	200			1450			1000		
Families	50			500			350		

Data source: Department of Education, National Center for Education Statistics, Early Childhood Longitudinal Study, Birth Cohort (ECLS-B), Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File. Sample size rounded to nearest 50 per NCES requirements.

3.2 Estimation Results

In Table 2 I present the estimation results for the model in equation (9). For estimation weights, I used wave 3 longitudinal weights for the primary caregiver survey for the single and married/cohabiting mothers samples, and wave 3 longitudinal weights for observations in the resident father survey for the married/cohabiting fathers sample.¹⁸

Using the estimates of Table 2, the structural parametrs can be retrieved as per equations (10) - (12). These estimated parameters are presented in Table 3 (the standard errors of these estimates are obtained by the delta method).

I consider the null hypothesis that two-parent families are using the same technology as one-parent families to invest in their children. This requires that $\alpha_1 = \alpha_2$, $\alpha_3 = 0$, and $\eta_1 = \eta_2$ (hypotheses 1, 2, and 3, respectively). The results of the test of this hypothesis are presented in Table 4. In that table, the number of observations used for the degrees of freedom is rounded to the nearest 50, as per NCES requirements.

¹⁶185% of the poverty line is the income threshold below which families are eligible for the National School Lunch Program and the Special Supplemental Nutrition Programs for Women, Infants, and Children.

¹⁷For comparison, a similar statistics computed from the US Census show even higher income disparities. Single mothers of children under 5 have a poverty rate of 41%, while for married couples this is 6% (author's calculations). Discrepancies are due to my sample selection procedure.

¹⁸See the appendix for an outline of the survey structure of the ECLS-B. There are several questionnaires, each with its own set of weights.

Table 2: Model Regression Coefficients for Skill Investment Technology Estimation (FE Model)

	Estimated Value
$\beta_{0,1}$	-1.957*** (0.0858)
$\beta_{0,2}$	-0.556*** (0.0858)
$\beta_{0,3}$	-0.984*** (0.0858)
β_1	-0.449*** (0.127)
β_2	-0.485*** (0.0639)
R^2	.5834
Observations	2650
SM families	50
MC families: mothers	500
MC families: fathers	350

Standard errors in parentheses. * p<0.05, ** p<0.01, *** p<0.001 Data source: Department of Education, National Center for Education Statistics, Early Childhood Longitudinal Study, Birth Cohort (ECLS-B), Longitudinal 9-Month-Kindergarten 2007 Restricted-Use Data File. Sample size rounded to nearest 50 per NCES requirements.

Table 3: Parameters of the Human Capital Accumulation Technology

η_1	α_1	η_2	α_2	α_3
0.449 (0.127)	0.0126 (0.0159)	0.485 (0.0639)	0.219 (0.0468)	0.0908 (0.0299)

Note: Standard errors are in parentheses, calculated for those parameters that are a function of estimated coefficients using the delta method.

Table 4: Hypothesis Test Results

	F	df	p-value
$H_{0,1}$	20.19	(1,1750)	0.0000
$H_{0,2}$	9.25	(1,1750)	0.0024
$H_{0,3}$	0.06	(1,1750)	0.7992
Joint:	406.99	(3,1750)	0.0000

Data source: U.S. Department of Education, National Center for Education Statistics, Early Childhood Longitudinal Study, Birth Cohort (ECLS-B), Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File. Sample size rounded to nearest 50 per NCES requirements.

These results establish three conclusions. First, as is evident from Table 4, one- and two-parent families do indeed use a different technology to invest in their children: the null hypothesis for the joint test is rejected at a 0.01% significance level. Second, these technologies have statistically the same elasticity of substitution between child care time and parental time. And third, the technologies have very different share parameters on child care time versus parental time.

The fact that the elasticities are not significantly different can be interpreted in terms of what that parameter governs in the investment technology. The estimation indicates that a 1% increase in the ratio of a single mother's hourly wage to the price of non-parental child care causes her to adjust the ratio of her time input to non-parental child care by 0.45%. She does not heavily readjust her investment input choices because of the price change. Married mothers, by comparison, adjust their input choices by 0.48% for a 1% change in p_t , very similar to single mothers. For married fathers, this statistic is by assumption the same as for the married mothers.¹⁹ These elasticities are informative about what one can expect parents in the economy to do with their time in the presence of a subsidy to the price of child care. All parents will shift inputs away from their own time and toward non-parental child care if the latter's price decreases. Ceteris paribus, this shift in input composition will not be dramatic; to achieve large changes in time use of the parents, the change in relative prices induced by a child care subsidy will have to be large.

The stark difference in the share on child care will drive differences in how families respond to a child care subsidy. Single mothers have a much higher share on the child care time input than couples do. In both family structures, however, the share on child care time is significant and larger than the shares on other inputs. If this share were instead zero, changes in the price of non-parental child care due to a subsidy would have no effect on the price of investment. In turn, this would mean that families did not change their choice of θ_{J+1} and skill accumulation in children would remain unaffected by the subsidy.

To summarize, if the price of child care decreases due to a subsidy, families will substitute towards child care in order to generate any given level of investment, and as they do so they will adjust the ratio of inputs

¹⁹Relaxing this assumption and allowing the elasticity to be different for mothers and fathers yields point estimates which are not statistically different from those presented here.

to a similar extent. They will also increase the target level of investment that they choose, because the price of each level of investment has gone down. Single mothers will increase their investment more than couples, *ceteris paribus*, because for single mothers the price of child care plays a larger role in determining the price of investment.

4 Heterogeneous Importance of Child Care Price

The price of child care—and therefore the effect of a child care subsidy—has different impacts across one- and two-parent families because these family types use different technologies to invest in their children, and because they face different income constraints. Different incomes arise because of the number of adults present in the household, and because of endogenous sorting by adult skill into family structures in the marriage market.

In order to gain some insight into the extent that each of these differences may determine differences in responsiveness to child care subsidies, in this section I derive an expression for the elasticity of skill's marginal cost with respect to the price of child care for each family structure. I decompose this into two components: a component that depends on the fraction of income being invested in the child, and the elasticity of the price of investment. Using the estimated parameters presented in the foregoing and ECLS-B data, as well as an assumption about the functional form of utility, I calculate and compare the empirical distributions of the marginal cost elasticity and its two components across family structures. Differences across family structures in two components will drive heterogeneity in response to the subsidy.

4.1 The Elasticity of the Marginal Cost of Investment With Respect To The Price of Child Care

From (1) and (2), consider a simple case where childhood lasts one period ($J = 1$), there is no borrowing or saving, and suppose the utility function in problems (1) and (2) takes the simple form $u(c, \ell) = \ln(c) + \psi^{type} \ln(\ell)$, where $type \in \{SM, MC\}$. Let (c_t^*, ℓ_t^*) denote the optimal solutions conditional on a given final skill for the child, θ_{J+1} . For single mothers this yields $c_t^* = \frac{w_{1,t} + T - \Lambda_t^{SM} I^{SM}}{1 + \psi^{type}}$ and $\ell_t^* = \frac{\psi^{type} (w_{1,t} + T - \Lambda_t^{SM} I^{SM})}{(1 + \psi^{type})(w_{1,t} + T)}$. Then the optimal choice of the final child skill, θ_{J+1} , maximizes $u(c^*, \ell^*) + bV(\theta_{J+1})$ and thus satisfies:

$$\frac{(1 + \psi)\Lambda_t^{SM}}{(w_{1,t} + T - \Lambda_t^{SM} I^{SM})} \frac{\partial I_t^{SM}}{\partial \theta_{J+1}} = b \frac{\partial V(\theta_{J+1})}{\partial \theta_{J+1}}$$

The left-hand-side of this equation displays the full “marginal cost” of child skill production in terms of utility, which balances the marginal benefit on the right-hand-side. Let MC^{SM} denote this marginal cost of single mothers in period t , and consider its elasticity with respect to the price of child care, $\rho^{SM} \equiv$

$\frac{\partial MC_t^{SM}}{\partial p_t} \frac{p_t}{MC_t^{SM}}$. From the foregoing:

$$\rho_t^{SM} = \left(\frac{w_{1,t} + T}{w_{1,t} + T - \Lambda_t I_t(\theta_{J+1})} \right) \epsilon_t^{SM} \quad (13)$$

where $\epsilon_{SM} = \frac{\partial \Lambda_t^{SM}}{\partial p_t} \frac{p_t}{\Lambda_t^{SM}}$ is the elasticity of the price of investment with respect to the price of child care. Given (5), this elasticity is:

$$\epsilon_t^{SM} = \frac{(1 - \alpha_1)^{\eta_1}}{(1 - \alpha_1)^{\eta_1} + \alpha_1^{\eta_1} \left[\left(\frac{w_{1,t}}{p_t} \right) \left(\frac{\phi_n}{\phi_1} \right) \right]^{1-\eta_1}} \quad (14)$$

Proceeding analogously for married couples, given (6) these objects are:

$$\rho_t^{MC} = \left(\frac{w_{2,t} + w_{3,t} + T}{w_{2,t} + w_{3,t} + T - \Lambda_t I_t^{MC}(\theta_{J+1})} \right) \epsilon_t^{MC} \quad (15)$$

$$\epsilon_t^{MC} = \frac{(1 - \alpha_2 - \alpha_3)^{\eta_2}}{(1 - \alpha_2 - \alpha_3)^{\eta_2} + \alpha_2^{\eta_2} \left[\left(\frac{w_{2,t}}{p_t} \right) \left(\frac{\phi_n}{\phi_2} \right) \right]^{1-\eta_2} + \alpha_3^{\eta_2} \left[\left(\frac{w_{3,t}}{p_t} \right) \left(\frac{\phi_n}{\phi_3} \right) \right]^{1-\eta_2}} \quad (16)$$

It is apparent that the higher the elasticity of investment price with respect to the price of child care, the more responsive the marginal cost will be to the child care subsidy, and the larger the expected increase in child skill, *ceteris paribus*. This elasticity is increasing in the share on child care time, and decreasing in the elasticity of substitution across inputs. Family income plays an additional role, however, as represented by the term multiplying the price elasticity to yield the investment elasticity. This income-composition term is increasing in the fraction of income spent on investment in children's skill.

4.2 Empirical Distributions of ρ_t^{SM} and ρ_t^{MC}

Table 5 reports the empirical distribution of the elasticities in (13) and (15), computed using the estimated parameters and ECLS-B data. Potential income is not measured (I only observe realized income), and so in order to calculate investment elasticities I use the estimated fixed effects at the parent level, which are necessarily not very precise.²⁰ Nevertheless, the table does yield two qualitative points. First, the mean and median of the elasticity of investment's price with respect to the price of child care are higher for single mothers than for couples, as anticipated from the point estimates for each family structure. Second, the fraction of income spent on investment in children's skill has a very similar mean and median for the two family structures. It is the empirical differences between ϵ_t^{SM} and ϵ_t^{MC} that drive the differences between the implied elasticity of the marginal cost of investment to the price of child care. Specifically, in the ECLS-

²⁰ Recall that each family fixed effect is identified based on only three observations. Single mothers in Table 5 are the same group as the sample. Couples in Table 5 are those families for whom both the father and mother are in the estimation sample. This is an additional restriction relative to the investment technology estimation. It is necessary to impose this restriction because the statistics reported here vary at the family level, not at the level of the parent.

B single mothers have on average a much larger marginal cost elasticity than couples do: it is almost three times higher for the former group.

Table 5: Empirical Elasticities: ρ_t^{type} and its Components

		mean	p50	sd
Single Mothers:	$\frac{income}{income - inv. expense}$	1.21	1.16	0.29
	ϵ_t^{SM}	0.87	0.89	0.09
	ρ_t^{SM}	1.06	1.02	0.31
	Observations	200		
	Families	50		
Couples:	$\frac{income}{income - inv. expense}$	1.29	1.25	0.20
	ϵ_t^{MC}	0.29	0.27	0.12
	ρ_t^{MC}	0.36	0.34	0.16
	Observations	900		
	Families	300		

Estimation Sample, Waves 1-3, Unweighted. Data source: U.S. Department of Education, National Center for Education Statistics, Early Childhood Longitudinal Study, Birth Cohort (ECLS-B), Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File. Sample size rounded to nearest 50 per NCES requirements.

In the following section, I develop a general equilibrium framework which I use for policy analysis. When I evaluate the model's fit using untargeted moments, I also compute an associated set of model statistics with their counterparts in Table 5. This comparison illustrates that the model population's sensitivity to the child care subsidy is in accordance with the actual population, when compared in terms of the statistics I derived here.

5 Investment in Child Skill in General Equilibrium

In this section I build a general equilibrium framework in to perform counterfactual policy analysis. In this expanded framework, the parenting problem for couples and single mothers during early childhood remains the same. The life cycle of the parents, however, now includes periods after children leave home. In addition, this framework endogenizes the family formation decision via marriage market. This necessitates specifying the problem of the single father to make his outside option to marriage concrete. To model the family formation decision, a small adjustment of notation is needed: the productivity of parents is now indexed by their gender, and refers to productivity in both the labor market and as a skill investment input. Accordingly, the symbol θ (with appropriate subscripts) is used to refer to this general productivity of the parent.

5.1 The Expanded Model

There are four sets of agents in the economy: consumers, a representative firm, the government, and a non-parental child care provider.

Given prices for labor and capital, the firm chooses labor and capital inputs to maximize profits. This firm produces with a Constant Returns to Scale technology and takes prices as given. The government chooses labor income taxes to finance lump-sum transfers and non-parental child care subsidies. A child support system exists, enforced by the government, where single fathers contribute a lump-sum amount that is redistributed lump-sum and equally to all single mothers. Finally, the non-parental child care sector supplies child care at the amount demanded in equilibrium. Consistent with the identifying assumptions in the estimation, I assume that the productivity of child care relative to parental time has a certain mean within one-parent families and two-parent families, respectively. To reflect that one- and two-parent families face similar relative prices, while single mothers also exhibit lower hourly wages, I allow the price of child care to differ by family structure.

The Life Cycle of Consumers Figure 1 illustrates the life cycle of the consumer. Each individual lives for $T + J$ periods. During childhood, which lasts for the first J periods of life, an individual makes no decisions: she is a passive recipient of consumption and investment chosen by her family. Upon independence, at the beginning of age $J + 1$, the individual leaves with the level of skill she has accumulated by then to start the remaining T periods of her life as an independent decision-making adult. As an adult, an individual operates as part of a family: families are either married/cohabiting couples (MC), single mothers (SM), and single fathers (SF). These families are formed in a marriage market that occurs at the start of period $J + 1$. The first J periods of adulthood are spent either actively parenting children (if a single mother or a married couple) or making child support payments (if a single father). From periods $2J + 1$ to T , the problem of the consumer is a standard lifecycle problem. As noted earlier, there are no wealth bequests—families make transfers to their children by increasing their skill.

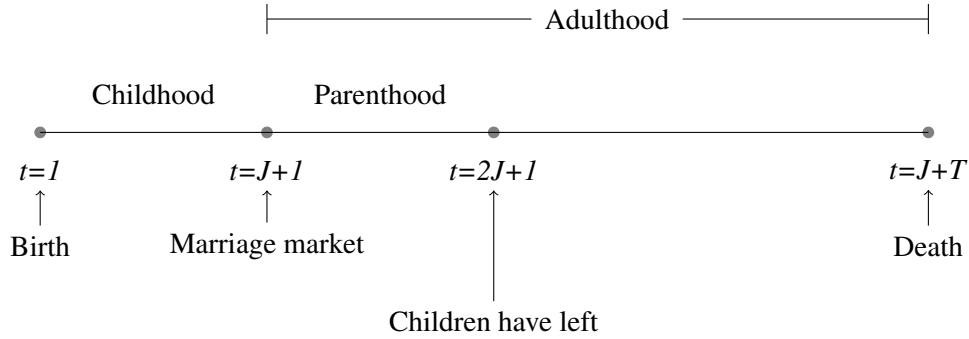


Figure 1: Life Cycle of the Consumer

Family Problems The family structures differ in three ways. First, they differ in their efficiencies of consumption: for single mothers and couples, these are denoted $\{\Phi_t^{SM}\}_{t=J+1}^T$ and $\{\Phi_t^{MC}\}_{t=J+1}^T$, which

represent consumption equivalence (CE) scales that vary over the lifecycle as children leave the household. The consumption equivalence scales of the single father are always equal to 1. Second, the different types of families are allowed to have different marginal utilities of leisure (which will be reflected in the parameters of the period utility functions (u^{SF}, u^{SM}, u^{MC})). Third, each family structure that raises children—single mothers and married/cohabiting couples—uses a skill accumulation technology specific to their family structure to invest in their child. This is consistent with the simpler framework used earlier in the paper, but with a shift of notation for the parenting productivities and hourly wage.²¹

The solution to a family’s life-cycle problem is a set of choices and lifetime utilities, for every potential couple $\{\theta_m, \theta_f\}$ and each possible draw of the initial child skill θ_1 . Here, θ_f and θ_m denote the female and male parent’s productivities, respectively. The lifetime utility for a single mother of type θ_f and child of type θ_1 is $V^{SM}(\theta_1, \theta_f)$, the lifetime utility for a single father of type θ_m with partner θ_f and child θ_1 is $V^{SF}(\theta_1, \theta_f, \theta_m)$, and the lifetime utility of a couple with types θ_f and θ_m , whose child has type θ_1 , is $V^{MC}(\theta_1, \theta_f, \theta_m)$.

The Marriage Decision Families are formed at the beginning of adulthood, when everyone participates in a marriage market that occurs instantaneously at start of period $J + 1$.²² On the marriage market, a potential match is drawn randomly from the skill distribution of the other gender in the same generation. Once assigned a potential spouse, and knowing that parenthood is certain in the environment, the agent compares the expected present discounted value of parenting alone or in a couple. A marriage is formed if both the husband and wife accept the match (the two individuals remain single otherwise). In particular, a female young adult with skill θ_f and potential spouse θ_m compares the expected value of being a single mother, $\mathbb{E}_{\theta_1} [V^{SM}(\theta_1, \theta_f)]$ and the expected value of being a married mother, $\mathbb{E}_{\theta_1} [V^{MC}(\theta_1, \theta_f, \theta_m)]$. A male young adult with skill θ_m and potential spouse θ_f compares the expected value of being a single father, $\mathbb{E}_{\theta_1} [V^{SF}(\theta_1, \theta_f, \theta_m)]$ and the expected value of being a married father, $\mathbb{E}_{\theta_1} [V^{MC}(\theta_1, \theta_f, \theta_m)]$.²³

After the marriage market, single mothers and couples draw the initial skill of their two children, which is the same for both children (single mothers and couples each raise both a son and a daughter). Whether parenting alone or in a couple, the lifetime utility of any individual contains a term that incorporates rational expectations about the lifetime utility of one’s child at the level of skill they begin adulthood with.²⁴ In the altruism term, the expected lifetime utility at adulthood, conditional on the level of skill, is taken over

²¹In particular, here the productivity of individuals as parents and in the labor market is assumed to be the same, and there is a common return per efficiency unit across individuals.

²²Appendix I contains transition matrices across waves in the ECLS-B for the mother’s marital status. In the data, it is extremely persistent.

²³The outside option to parenting in a couple differs by gender in the model because, empirically, the vast majority of single parents who are raising young children in their home are women. For a discussion of what the ECLS-B offers in terms of discipline on contributions of parental time from single fathers, see Appendix G.

²⁴This is what makes parents altruistic. An alternative way of motivating intergenerational transfers is through paternalistic preferences, or “warm glow” returns (Andreoni (1990)). The advantage of an altruism framework is that the returns to investment can respond endogenously to policy, because parents fully incorporate the economic returns to their investment in terms of their child’s lifetime utility, and their behavior changes accordingly. The main benefit of a paternalistic specification is its tractability and flexibility in matching parenting behaviors. Some models include both altruism and a paternalistic preference for, say, college attainment which is distinct from its monetary returns. For an application of paternalistic preferences to intergenerational transfers of wealth, see De Nardi (2004).

the potential spouses one might meet as an adult (using the distribution of skill in the economy, which is endogenous) and also over the distribution of initial skill one's child may be born with (which is exogenous).

The decision rule that maps from the type of the spouse to “no” or “yes” marriage market decision is represented by $D_g(\theta_m, \theta_f) \in \{0, 1\}$. It takes as given the skill θ_m or θ_f and gender g of the decision maker, which can be either male, m , or female, f , and solves:

$$D_f(\theta_m, \theta_f) = \arg \max_{\delta \in \{0,1\}} \left[\delta \int_{\theta_1} V^{MC}(\theta_1, \theta_f, \theta_m) \pi(\theta_1) d\theta_1 + (1 - \delta) \int_{\theta_1} V^{SF}(\theta_1, \theta_f, \theta_m) \pi(\theta_1) d\theta_1 \right] \quad (17)$$

$$D_m(\theta_m, \theta_f) = \arg \max_{\delta \in \{0,1\}} \left[\delta \int_{\theta_1} V^{MC}(\theta_1, \theta_f, \theta_m) \pi(\theta_1) d\theta_1 + (1 - \delta) \int_{\theta_1} V^{SM}(\theta_1, \theta_f) \pi(\theta_1) d\theta_1 \right] \quad (18)$$

where the distribution of initial child skill θ_1 has a probability density function given by $\pi(\theta_1)$. The optimal marriage decision takes the form of a threshold strategy in the space of the potential spouse's skill. The value functions that young adults compare when deciding whether to marry depend on the return to skill in terms of lifetime utility, which is constructed next.

Expected Lifetime Utility The expected lifetime utility of the family's two children (one boy and one girl) enters into the parent problem. It is:

$$V^{child}(\theta) = \sum_{g \in \{m,f\}} \tilde{V}_g(\theta) \quad (19)$$

This function sums the expected lifetime utility given skill θ across potential spouses, potential children, and the gender of the child. That is, each family internalizes the average return to skill in terms of lifetime utility across men and women in equilibrium when making investment decisions. For each gender the expected return to skill in terms of lifetime utility is:

$$\tilde{V}_f(\theta) = \int_{\theta_m} \left[\mathbb{I}_d \int_{\theta_1} V^{MC}(\theta_1, \theta, \theta_m) \pi(\theta_1) d\theta_1 + (1 - \mathbb{I}_d) \int_{\theta_1} V^{SM}(\theta_1, \theta) \pi(\theta_1) d\theta_1 \right] \mu(\theta_m) d\theta_m \quad (20)$$

$$\tilde{V}_m(\theta) = \int_{\theta_f} \left[\mathbb{I}_d \int_{\theta_1} V^{MC}(\theta_1, \theta_f, \theta) \pi(\theta_1) d\theta_1 + (1 - \mathbb{I}_d) \int_{\theta_1} V^{SF}(\theta_1, \theta_f, \theta) \pi(\theta_1) d\theta_1 \right] \mu(\theta_f) d\theta_f \quad (21)$$

where $\mathbb{I}_d \equiv D_f(\theta_m, \theta_f) \times D_m(\theta_m, \theta_f)$ indicates a mutual acceptance of the match. In addition, $\mu(\theta)$ is the endogenous distribution over adult skill θ , which is the same for each gender because parents are constrained to not targeting investments by gender.

Once families are formed, the family solves a sequential life-cycle problem. Just as in problems (1) and (2), the family has J periods to invest in their child. The child is born with some initial skill θ_1 . At each age during adulthood, t , the parent affects the child's stock of skill in that period, θ_t , by their choice of investment, I_t^{type} , where $type \in \{SM, MC\}$ for single mothers and married couples, respectively. Note

that age t of adulthood is age $t + J$ of the parents' life. The way investment interacts with the child's skill in each period is defined by the skill accumulation function $f(\theta_t, I_t^{type})$. In turn, the way investment in each period is generated depends on the skill investment technology, which is indexed by family structure. Investment is generated by time contributed from the parent(s) (q_t^m and q_t^f) and from time purchased on the market in the form of non-parental child care n_t .²⁵ In all the family problems, w is the return per unit of skill on the labor market, r is the interest rate earned from the stock of savings, τ_y is the labor income tax, τ_n is the subsidies to child care, T are lump-sum transfers, and T_{cs} are child support payments.

Single Mothers The problem of a single mother, articulated in (1) but with the aforementioned updated notation, can now be stated as:

$$\begin{aligned}
V^{SM}(\theta_1, \theta_f) &= \max_{\{c_t, \ell_t, a_{t+1}, n_t, q_t^f\}_{t=J+1}^T} \left[\sum_{t=1}^T \beta^{t-1} u^{SM} \left(\frac{c_t}{\phi_t^{SM}}, \ell_t \right) \right] + \beta^{J-1} b V^{child}(\theta_{J+1}) \quad (22) \\
s.t. \\
c_t + a_{t+1} + (1 - \tau_n) p_t^{SM} n_t &\leq w \theta_f (1 - \ell_t - q_t^f) + (1 + r) a_t + T + \mathbb{I}_{t \leq J} T_{cs} \\
\theta_{t+1} &= f(\theta_t, I_t^{SM}) \\
I_t^{SM} &= I_t^{SM}(\phi_n^{SM} n_t, \theta_f q_t^f)
\end{aligned}$$

where, as before, $a_1, a_{T+1} = 0$, $n_t, \ell_t, q_{1,t} \geq 0$, and the time constraints are $\ell_t + q_{1,t} \leq 1$ and $n_t + q_{1,t} \leq 1$ (both parents and children have one unit of time per period). Here, b is the altruism parameter, $w \theta_f$ denotes the wage of the parent in period t , and p_t^{SM} denotes the price of non-parental child care in period t . As before, parental and non-parental child care time investments in child skill affect skill in the next period according to the skill production function $f(\theta_t, I_t^{SM})$ and the investment function $I_t^{SM}(\phi_n^{SM} n_t, \theta_f q_t^f)$. Note that the investment function I_t^{SM} is indexed to the family structure, while the production function f is the same for both types of families. The final child skill θ_{J+1} enters the objective function of the mother through an altruism term $b V^{child}(\theta_{J+1})$, which weights the expected lifetime utility of the child $V^{child}(\theta_{J+1})$ with the altruism coefficient b .

Married or Cohabiting Couples Similarly, the cohabiting (which includes married) couples problem in (2) in the expanded framework's notation is stated as:

$$\begin{aligned}
V^{MC}(\theta_1, \theta_f, \theta_m) &= \max_{\{c_t, \ell_t, a_{t+1}, n_t, q_t^m, q_t^f\}_{t=J+1}^T} \left[\sum_{t=1}^T \beta^{t-1} u^{MC} \left(\frac{c_t}{\phi_t^{MC}}, \ell_t \right) \right] + \beta^{J-1} b V^{child}(\theta_{J+1}) \quad (23) \\
s.t. \\
c_t + a_{t+1} + (1 - \tau_n) p_t^{MC} n_t &\leq w \theta_f (1 - \ell_t - q_t^f) + w \theta_m (1 - \ell_t - q_t^m) + (1 + r) a_t + T \\
\theta_{t+1} &= f(\theta_t, I_t^{MC}) \\
I_t^{MC} &= I_t^{MC}(\phi_n^{MC} n_t, \theta_f q_t^f, \theta_m q_t^m)
\end{aligned}$$

²⁵Recall that the superscripts m and f now refer to the male and female parent.

where, as before, $a_1, a_{T+1} = 0$, $n_t, \ell_t, q_{2,t}, q_{3,t} \geq 0$, and the time constraints are $\ell_t + q_{2,t} \leq 1$, $\ell_t + q_{3,t} \leq 1$, and $n_t + \max\{q_{2,t}, q_{3,t}\} \leq 1$. Parents are allowed to invest in the child at the same time (their quality time investments are non-rival). More generally, the specification used here for the married couple problem is based on Guvenen and Rendall (2015). As in that study, the perfect complementarity in leisure of the spouses is motivated with time use data as documented in Aguiar and Hurst (2007), Table V. This is reflected in the parenting problem (2) by the fact that the couple only chooses one level of leisure which they both enjoy.²⁶

Single Fathers A single father chooses consumption c_t , savings a_{t+1} , and leisure ℓ_t to solve the following problem.

$$V^{SF}(\theta_1, \theta_f, \theta_m) = \max_{\{c_t, \ell_t, a_{t+1}\}_{t=J+1}^T} \left[\sum_{t=1}^T \beta^{t-1} u^{SF}(c_t, \ell_t) \right] + \beta^{J-1} bV^{child}(\mathbb{E}(\theta_{J+1}|\theta_1, \theta_f)) \quad (24)$$

$$c_t + a_{t+1} + \mathbb{I}_{t \leq J} T_{cs} \leq w\theta_m(1 - \ell_t) + (1 + r)a_t + T$$

where $\ell \in [0, 1]$, and $a_1, a_{T+1} = 0$. Here $w\theta_m$ is wage of the single father with skill θ_m . Single fathers take their child's outcome θ_{J+1} as given conditional on the female parent's type, θ_f , as well as the child's initial skill θ_1 . A single father cannot use his own time to invest in his children, but he is required to make a lump-sum payment to the child's mother, T_{cs} . The expected outcome of a single father's children affects his outcome directly through the altruism term, $bV^{child}(\mathbb{E}(\theta_{J+1}|\theta_1, \theta_f))$.

Government The government collects revenue from labor income taxes τ_y to finance lump-sum transfers T and non-parental child care subsidies τ_n . The variable H_t is the aggregate supply of labor efficiency units at each age during adulthood t , N_t^{type} is the aggregate demand for non-parental child care from parents in family structure $type \in \{SM, MC\}$.

$$\tau_y w \sum_{t=J+1}^T H_t = T + \tau_n \sum_{type} p_n^{type} \sum_{t=J+1}^{J+J} N_t^{type} \quad (25)$$

Representative Firm The firm chooses capital K_F and labor inputs H_F to maximize profits, taking prices r and w as given. The parameter δ_F is the depreciation rate of capital.

$$\max_{K_F, H_F} \left\{ K_F^{\alpha_F} H_F^{1-\alpha_F} - wH_F - (r + \delta_F) K_F \right\} \quad (26)$$

Non-parental Care Sector The non-parental child care sector provides N units of non-parental child care at price p_n^{type} . The price of non-parental child care is set as a constant fraction κ^{type} of the average return to labor supply for female parents in family structure $type \in \{SM, MC\}$.

²⁶Unlike Guvenen and Rendall (2015), here the marginal utility of leisure is deterministic.

$$p_n^{type} = \kappa^{type} (1 - \tau_y) w \int_{\theta_f} \theta_f \mathbb{I}_{type} \mu(\theta_f) d\theta_f \quad (27)$$

This allows the price of non-parental child care to adjust with the average level of skill in the economy, but without specifying a production function for non-parental child care.²⁷

Equilibrium Given a government policy τ_n , transfers T , and child support T_{cs} , a stationary equilibrium is defined as factor prices and a labor income tax, individual marriage decisions, family decisions and lifetime utilities, expected lifetime utilities at a given level of adult skill, and single father expectations about single mother decisions such that markets clear, the government balances its budget constraint, consumers and firms optimize and expectations are rational.

5.2 Parameterization

In order to fully parameterize the general equilibrium model, I assume functional forms for utility, the initial distribution of child skill and its correlation with parental attributes, and $f(\theta_t, I_t^{type})$, the function that governs how investment and skill today combine to produce tomorrow's skill. I assign parameters to these functions via calibration and estimation, respectively.

Period Utility Functions Utility functions are defined separately for one- and two-parent families:

$$\begin{aligned} u^{SM}(c, \ell) &= \log(c) + \psi_s \log(\ell) \\ u^{SF}(c, \ell) &= \log(c) + \psi_s \log(\ell) \\ u^{MC}(c, \ell) &= \log(c) + \psi_{mc} \log(\ell) \end{aligned}$$

Distribution of Initial Child Skill I assume that initial child skill is drawn independently and identically from $\pi(\theta_1)$, which I set as a uniform distribution. The empirical motivation for the i.i.d. assumption is based on the correlation between the cognitive score at 9 months of age and after-tax family income in the ECLS-B, which is 0.01 with a p-value of 0.12. This correlation is not statistically different from 0.

Skill Production Function The child's current stock of skill, θ_t , and chosen level of investment I_t^{type} combine according to the skill production function $f(\theta_t, I_t^{type})$ to produce the stock of skill in the next period. For both family structures, I assume the same CES functional form:

$$\theta_{t+1} = f(\theta_t, I_t^{type}) = \left[v \left(\lambda_{type} I_t^{type} \right)^{\frac{\chi-1}{\chi}} + (1-v) (\theta_t)^{\frac{\chi-1}{\chi}} \right]^{\frac{\chi}{\chi-1}}$$

In a one-period parenting problem parameter, the elasticity of substitution parameter χ can be interpreted

²⁷Because parents cannot target investment to the child based on their gender, the skill distribution is the same for both genders. That means that, within the model, using the skill distribution of fathers to pin down the price of child care would be equivalent.

as governing how easily investment can compensate for initial skill. If this elasticity is low, it takes more investment to change initial skill than if the elasticity is high. In a multi-period setting where $J > 1$, the elasticity of substitution χ also governs how easily investment can be shifted across periods of childhood in response to changes in its price across periods. If χ is low, the technology does not allow investment to be easily reallocated across periods in response to a change in the price of investment in one period. In a one-period parenting problem, the share parameter v can be interpreted as determining the importance of investment in determining the final skill of the child. If v is low, it takes more investment to affect the stock of skill, and vice versa. In a model where $J > 1$, the share parameter v also determines how investment with a constant price across consecutive periods would be allocated. If v is low, then investment in each period of childhood will have to be more uniform, whereas if v is high then investment could be more lumpy. Finally, λ_{type} is a scaling parameter that accounts for units of measurement being different across investment and the stock of skill (this parameter is allowed to differ across family structures).

Estimating χ and v To estimate the parameters of the production function $f(\theta_t, I_t^{type})$, I rely on the tangency conditions from inter-temporal cost minimization. First, note that the child's final skill depends on the initial condition and the sequence of investments chosen by the family. The implied function governing the mapping from initial to final skill can be constructed recursively as:

$$\bar{\theta}_{J+1} \left(\theta_1, \left\{ I_t^{type} \right\}_{t=1}^J \right) = \left[\sum_{t=1}^J v (1-v)^{J-t} \left(\lambda I_t^{type} \right)^{\frac{\chi-1}{\chi}} + (1-v)^J (\theta_1)^{\frac{\chi-1}{\chi}} \right]^{\frac{\chi}{\chi-1}}$$

Next, by the (conditional) cost minimization implicit in the optimization problems (22) and (23), one can derive an inter-temporal Euler equation which balances the relative price of investment in consecutive periods with the relative price of moving value across those periods. Specifically:

$$\frac{1}{1+r_{t+1}} \frac{\Lambda_t}{\Lambda_{t+1}} = \frac{\frac{\partial \bar{\theta}_{J+1}(\theta_1, \{I_t^{type}\}_{t=1}^J)}{\partial I_t}}{\frac{\partial \bar{\theta}_{J+1}(\theta_1, \{I_t^{type}\}_{t=1}^J)}{\partial I_{t+1}^{type}}} \quad (28)$$

Substituting the partial derivatives of $\bar{\theta}_{J+1}$ into (28) and rearranging yields the following estimation equation:

$$\ln \left(\frac{I_{i,t+1}^{type}}{I_{i,t}^{type}} \right) = -\chi \ln \left(\frac{\Lambda_{i,t+1}^{type}}{\Lambda_{i,t}^{type} (1+r_{t+1})} \right) + \chi \ln (1-v) \quad (29)$$

where the parameters of the unit cost functions $\Lambda_{i,t}^{type}$ were estimated earlier in the paper. This is the structural portion of a linear regression of the form $y_{i,t} = \gamma_0 + \gamma_1 x_{i,t} + e_{i,t}$, which I estimate by OLS. The last two skill production technology parameters can then be retrieved as $\chi = -\gamma_1$ and $v = 1 - \exp \left(-\frac{\gamma_0}{\gamma_1} \right)$. The results of this estimation are reported in Table 6, yielding a value for χ of 0.23 and a value for v of 0.27.

Table 6: Model Regression Coefficients

	Estimated Value
γ_1	-0.227* (0.0956)
γ_0	0.0723* (0.0336)
R^2	.0189
Obs.	700
SM families	50
MC families	300

Standard errors in parentheses. *
 $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ Data
source: U.S. Department of Educa-
tion, National Center for Education
Statistics, Early Childhood Longitu-
dinal Study, Birth Cohort (ECLS-B),
Longitudinal 9- Month-Kindergarten
2007 Restricted-Use Data File. Sam-
ple size rounded to nearest 50 per
NCES requirements.

Remaining Parameters The remaining parameters of this model are divided into three groups and presented in the panels of Table 7: fixed parameters in Panel A, estimated parameters in Panel B, and calibrated parameters in Panel C.

Beginning with Panel A, the length of a lifetime and of each phase are proportional to 20 years of childhood and 60 years of adulthood (death at age 80). The discount (patience) factor is set to a yearly value of 0.96 to match the presumed risk-free interest rate. The time endowment for early childhood is 5 years, which is 0.25 the length of a period (20 years). The share on capital, α_F , is set to 0, thereby preserving the assumption of constant returns to scale in production. Because of the long length of a period in this model, effectively this shuts off long-term borrowing for parents. The depreciation rate of capital, δ_F is set to 0. Finally, the CE scales are set using the 1994 scales from the Organisation for Economic Co-operation and Development (OECD). These scales assign a value of 1 for the first adult, and 0.5 for the subsequent adults; for each dependent the weight is 0.3. They adjust money spent on consumption into units of consumption for each member of the household. Once children leave the family, the CE scale for single mothers goes back to 1, and the scale for couples falls to 1.5.

Table 7: Fixed, Estimated, and Calibrated Parameter Values

	Symbol	Name	Source	Value
A. Fixed:	β	Patience	Assumption	0.96^{20}
	J	Duration early childhood	$\frac{20 \text{ year period}}{5 \text{ year early childhood}}$	0.25
	$\{\alpha_F, \delta_F\}$	Production technology	Assumption	$\{0, 0\}$
	$\{\phi_1^{SM}\}$	CE Scales: 1 ad., 2 ch.	OECD	$\{1.6\}$
	$\{\phi_1^{MC}, \phi_{2,3}^{MC}\}$	CE Scales: 2 ad., 2 ch.	OECD	$\{2.1, 1.5\}$
B. Estimated:	T	Transfers	NIPA	8% of output
	T_{cs}	Child Support	Census, NIPA	45 % of T
	κ^{type}	p_n^{type} coefficient	ECLS-B	35% \bar{w}_{mother}^{SM} , 39% \bar{w}_{mother}^{MC}
C. Calibrated:	b	Altruism coefficient		$0.78 \times \beta$
	ψ_s, ψ_{mc}	MU_ℓ singles and couples		2.4, 1.8
	$\lambda_{SM}, \lambda_{MC}$	Prod. of Inv. (SM,MC)		12, 25

Moving to Panel B, the level of lump-sum transfers T , the level of child support payments T_{cs} , and the price of p_n are set to 8% of output, 45% of the average per-family transfer, and 35% (single mothers) or 39% (married/cohabiting mothers) of the average mother's wage, respectively.. The first empirical target is from the ratio of government transfers to persons for federal benefits from social insurance funds, Supplemental Nutrition Assistance Program (SNAP), supplemental security income, refundable tax credits, and other (which includes payments to nonprofit institutions and student loans, among other categories) to GDP from the National Income and Production Accounts (NIPA) tabulations. The second is the ratio of average child support payments owed per month per capita to average monthly government transfers per family. The third is the average ratio of hourly price of non-parental child care to hourly wages of mothers in the ECLS-B. See the appendix for further details on the targets for these parameters.

Panel C presents the values for internally calibrated parameters —the coefficients $b, \psi_s, \psi_{mc}, \lambda_{SM}$ and λ_{MC} are chosen to bring the model moments in the baseline as close as possible to the moments in Table 8. The coefficient b controls the degree of altruism; ψ_s, ψ_{MC} are the marginal utilities of leisure for singles and married couples, respectively. The parameters λ_{SM} and λ_{MC} are shifters in the skill technology that scales up investment into efficiency units in the production of skill. The moments these parameters were chosen to match, and the fit of the calibration, are reported in Table 8.

Table 8: Calibration Targeted Moments

Moment	Source	Data	Model
$Corr(\text{Skill}_{\text{child}}, \text{Inc.}_{\text{family}}) \text{ age } 5$	ECLS-B	0.32	0.32
Ave. labor supply	CPS	0.31	0.31
Pct. families MC	ECLS-B	0.81	0.81
Parental time: MC ($\frac{\text{Hours}}{\text{Week}}$)	ECLS-B	(9.3, 4.8)	(9.4, 8.0) (MM, MF)
Parental time: SM ($\frac{\text{Hours}}{\text{Week}}$)	ECLS-B	2.5	4.7 SM

The moments in Table 8 are the correlation of child skill with family income, the average labor supply of parents with children under 5 who are between 15 and 55 years old, the percent of single mothers raising children under 5, and the average time invested by parents of each family structure type. The target moments from the ECLS-B are from the pooled estimation sample whose moments are given in Table 1, except for the marriage rate, which is from the unrestricted weighted sample described in 12 in the appendix.²⁸

The moment most poorly matched in this calibration is the hours of parental time provided by single mothers and married/cohabiting fathers, which are too high. This is driven by the fact that λ_{SM} has to be adjusted downward relative to λ_{MC} to match the marriage rate in equilibrium.

Untargeted Moments Table 9 compares the equilibrium moments of this model with untargeted moments in the data. These moments fall into two broad categories: family formation and inputs, presented in Table 9a, and investment elasticities, presented in Table 9b.

Table 9a shows that the random search marriage market captures a quantitatively reasonable degree of assortative matching among spouses: the correlation of wages within a couple is slightly lower than the data at 0.22 versus 0.27. The lower equilibrium value could be due to unmodelled types which couples sort by in the data but not the model, such as college attainment, or to an unmodeled segmentation of the marriage market. The second moment shows that the correlation of time inputs in the model and the data are quite close, although again the model demonstrates a slightly lower value for the statistic at 0.35 instead of 0.39. Finally, the ratio of time input levels from mothers and child care is again slightly lower (0.33) than the data (0.36). Overall, the model is parsimonious but generates correlation in parental types and investment input choices which align with their empirical counterparts.

Table 9b reports another set of statistics by which I compare the baseline equilibrium with the data: the elasticities of the investment price and the marginal cost of skill. These statistics are important for the responsiveness of the model to a child care subsidy, and they are untargeted in the model calibration. Comparing the model with the data, the elasticities are quite close for single mothers but slightly higher for

²⁸Note that the correlation of child skill and family income grows to 0.32 by the time the child is 5 years old; by contrast, when children are 9 months old, the measures of skill available in the data are uncorrelated with family income, although they do have predictive power for later child test scores at age 4. For regressions supporting these points, especially the assumption that initial child skill is independent of family attributes, see Appendix II.

couples. This is because of the higher investment elasticity in the model versus the data for couples, which in turn is a result of stronger sorting by parent skill into marital status than is present empirically. In the next section, I examine the relationship between family income, family structure, and child outcomes more closely.

Table 9: Untargeted Moments

(a) Family Formation and Input Choices

Moment	Data	Model
Correlation of wages within a couple	0.27	0.22
Correlation of time inputs within a couple	0.39	0.35
Average $\frac{\text{mother time}}{\text{non-parental child care time}}$	0.36	0.33

(b) Elasticity of Marginal Cost of Investment ρ_t^{type} and its Components

	Moment (means)	Data	Model
Single Mothers:	$\frac{\text{income}}{\text{income} - \text{inv. expense}}$	1.21	1.28
	ϵ_t^{SM}	0.87	0.86
	ρ_t^{SM}	1.06	1.10
Couples:	$\frac{\text{income}}{\text{income} - \text{inv. expense}}$	1.29	1.10
	ϵ_t^{MC}	0.29	0.39
	ρ_t^{MC}	0.36	0.43

Note: Table 9a correlations are Pearson correlation coefficients, weighted with survey weights in the ECLS-B and using the analogous distribution in the model. Note: Data moments in Table 9b are from Table 5; model moments are computed from model output.

6 Policy Experiment: Universal Child Care Subsidy

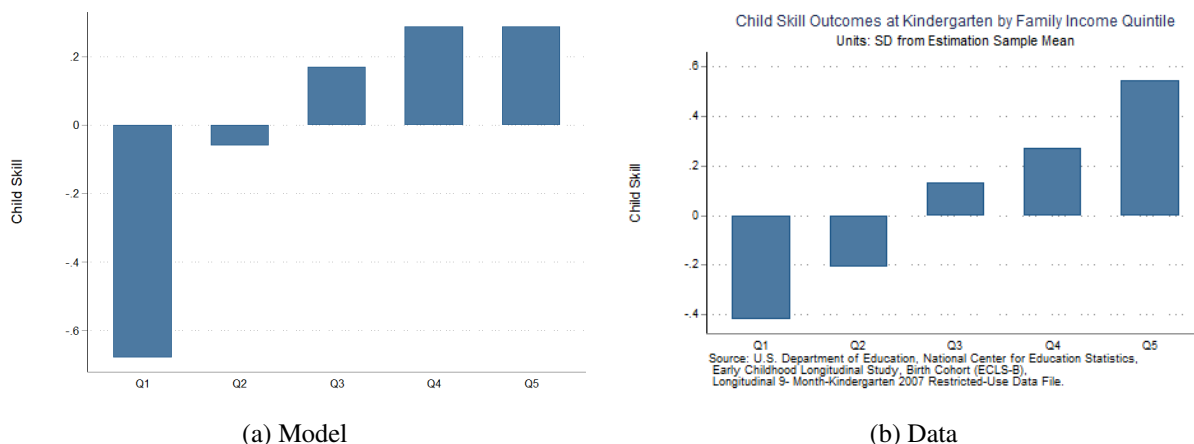
Before proceeding to analyzing the effects of universal child care subsidies in general equilibrium, I review how family income, family structure, and child outcomes are related in the baseline equilibrium of the model. In my analysis of the policy experiment, I return to these statistics and evaluate how the subsidy affects them.

6.1 Family Income, Family Structure, and Child Outcomes

In this model, initial child skill is independent of family attributes, but child outcomes are not. The correlation of child skill with family income is a targeted moment in the calibration of the model, which generates a pattern of child skill outcomes being increasing by family income quintiles (Figure 2a). The relationship

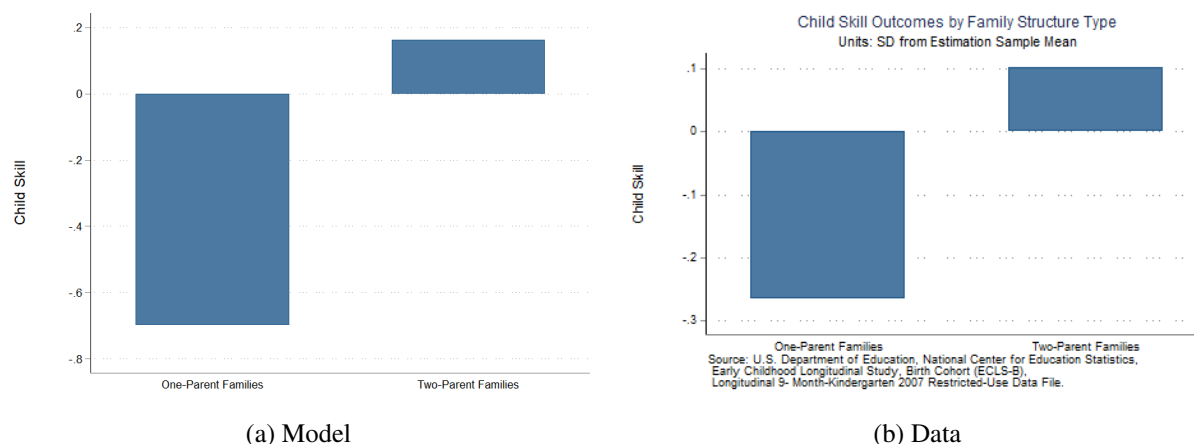
between child skill outcomes and family structure, on the other hand, is not targeted in the calibration and is an equilibrium result. Figure 3a presents the average child skill in equilibrium by family structure: children of single mothers are on average about 0.9 standard deviations lower than the average outcome for couples.

Figure 2: Child Skill Outcomes and Family Income: Model and Empirical Moments



This is qualitatively consistent with the empirical patterns presented in Figures 2b and 3b, which use information on child skills and family attributes measured in the ECLS-B.²⁹ Empirically, when children are 9 months old, measures of skill and after-tax family income exhibit a correlation coefficient of 0.01 (with a p-value of 0.12). When they are 4 years old, however, this correlation jumps to 0.32 and is significant at the 0.01 level.

Figure 3: Child Skill Outcomes and Family Structure: Model and Empirical Moments




Finally, family structure is correlated with income both in the model and in the data. In Figure 4a, the marriage rate within each income quintile is shown to be strongly increasing in income. This is qualitatively

²⁹See Appendix I for a description of skill assessments for children in the ECLS-B. Measures of child skill, using age-appropriate assessments, were collected in each wave of the survey.

consistent with what is observed in the ECLS-B data (Figure 4b). However, the model exhibits income differences across family structures which are more extreme than what is present in the data.

Figure 4: Marriage Rate by Family Income Quintile: Model and Empirical Moments



6.2 Universal Child Care Subsidies

With the parameterized model in hand, I vary the child care subsidy τ_n between 0 and 0.99, solve for the stationary equilibrium, and record welfare, child skill outcomes, family income, and choices of child care and parental quality time. I also record labor income taxes that balance the government's budget constraint at every level of the child care subsidy. Because marriage is endogenous, the model keeps track of family decision and child outcome variables conditional on the marital status, so that I can construct the counterfactual outcome for parents and children if the couple had formed (or not) instead of what was optimally chosen in equilibrium.

This model exhibits heterogeneity in the effects of child care subsidies across family income and family structures. The model also proposes a mechanism by which family income and family structure are correlated endogenously via the marriage market decision. Using the model framework to observe a counterfactual scenario where sorting into family structure is random, I decompose the difference in child outcomes across family structures into a component due to the family structure technology, and a portion due to the compositional differences across family types. As the gap narrows due to the child care subsidy, the role of technology in driving outcome differences is computed and compared over different levels of the subsidy.

6.2.1 Child Care Subsidies and Child Skill Accumulation in Equilibrium

Using the model output, I document several properties of the child care subsidies effect on this environment. First, the subsidy narrows the gap in child outcomes between one- and two-parent families. This occurs because the subsidy leads to a larger increase in the outcomes of children from one-parent families relative to two-parent families, due to both the lower income of single mothers and the technology they use to invest in their children. Second, the subsidy does not strongly affect family formation decisions or the family

structure composition of different income quintiles, and it leads to an increase in the labor income tax in order to fund it, which lowers the return to skill in the market and decreases the marginal benefit of skill. This force offsets the decrease in the marginal cost of skill due to the subsidy, and eventually dominates it. Thus, at a very large subsidy, average child outcomes are decreasing in the level of the subsidy. Third, welfare gains are maximized where gains in child skill reach their highest level.

Levels of child skill, in units of standard deviations from the population mean, are presented in Figure 5a. These gains are due to the subsidy lowering the cost of investment. As a result, children do better, and become more productive adults, both as workers and as parents. These effects are stronger for one-parent families, which is in line with the attributes of their investment technology as summarized by the elasticity of marginal cost with respect to the price of child care. It is also in line with the findings of Kottelenberg and Lehrer (2017) and Havnes and Mogstad (2014), who document higher gains for children of one-parent, poor households. Because of this difference in child skill gains, the gap in child outcomes across one- and two-parent families narrows as the subsidy decreases (Figure 5b).

Figure 5: Heterogeneity in Effects of Child Care Subsidy Decreases the Gap in Child Outcomes

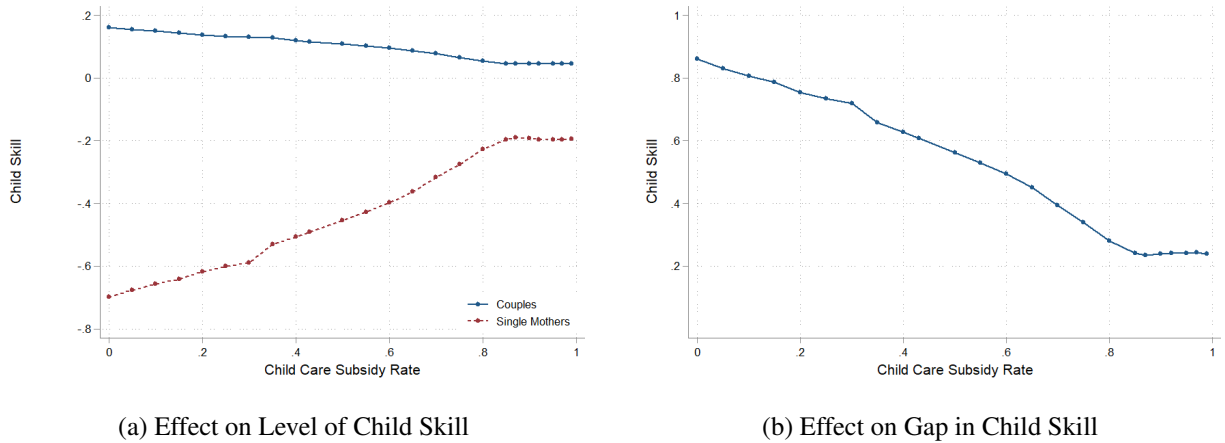
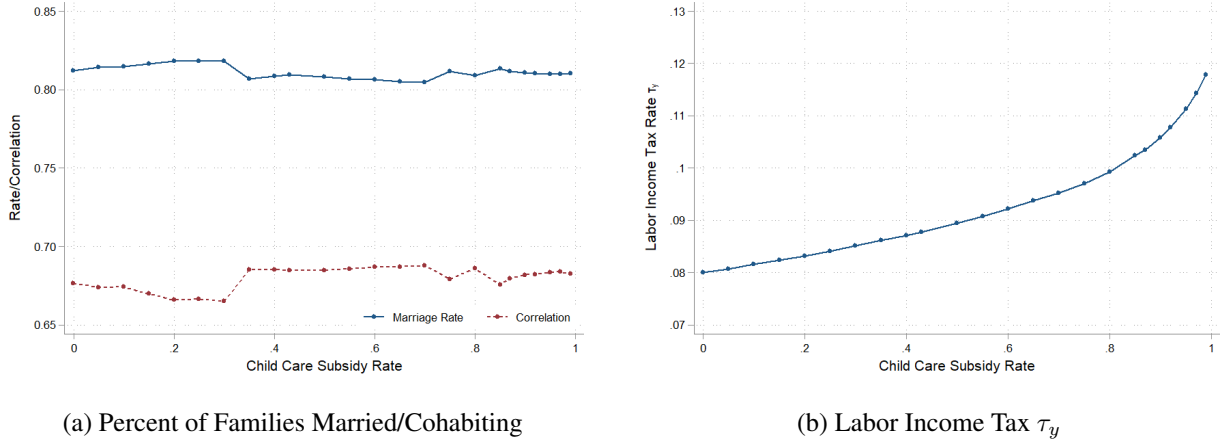


Figure 6a presents the percent of families who are married/cohabiting as well as the correlation between marital status and family income quintile, at each level of the child care subsidy. Neither statistic is strongly affected by the subsidy. In this model, endogenous family formation will matter in determining the skill composition of single mothers versus couples. However, the marriage decision is not a margin of adjustment in the face of changes to the price of child care, and the distribution of family structures over income quintiles is correspondingly stable.

As the child care subsidy increases, labor income taxes also increase to fund it. This is shown in Figure 6b. Eventually, the resulting losses in the return to skill on the market lead aggregate welfare to start decreasing in the child care subsidy.

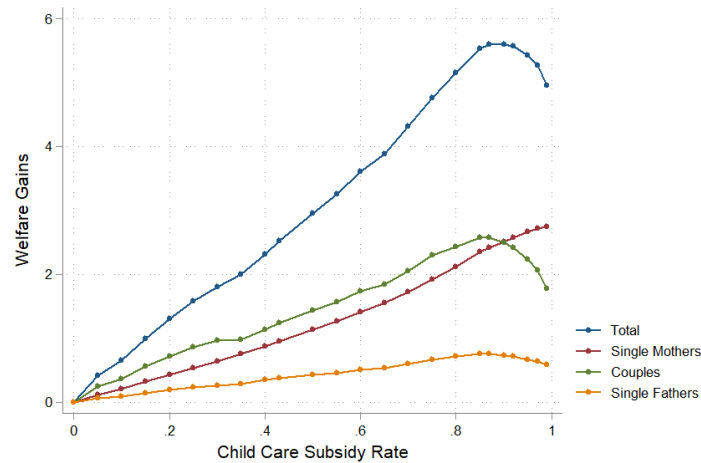
As the child care subsidy level increases, thereby decreasing the marginal cost of investment, the labor income tax increases as well in order to fund the subsidy. This increase in the labor income tax decreases the return to skill, $(1 - \tau_y)w$, so that the marginal benefit of skill in terms of lifetime utility is also decreasing

Figure 6: Effect of a Child Care Subsidy on Family Formation and the Labor Income Tax



in the subsidy. It is possible for these changes to offset one another, and they are more likely to do so for couples and single fathers at lower levels of the subsidy. This is because the price of investment is less affected by the price of child care for couples than for single mothers, due to the lower share on child care in their investment technology. Single fathers benefit vicariously from increases in child outcomes, but do not have to contribute from their budget constraint to finance child care of parental time. Instead, their contributions are lump-sum in the form of child support. Figure 7 shows the total welfare gains. It also reports additive weighted contributions from couples, single mothers, and single fathers to the total. It is evident that the three family types in the economy experience different patterns of welfare gains due to the subsidy.

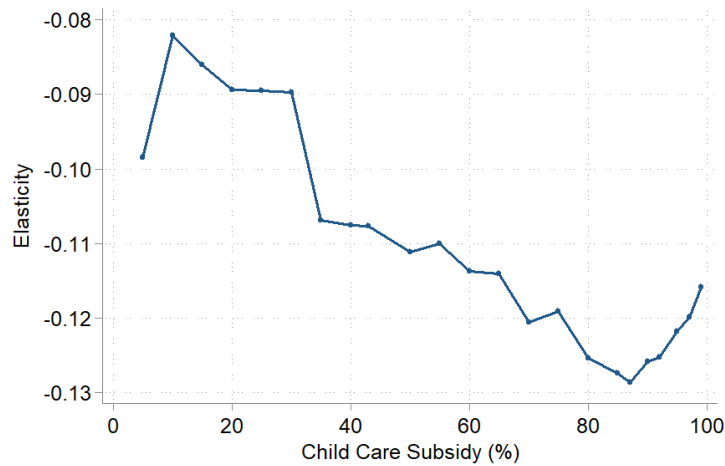
Figure 7: Effect of a Child Care Subsidy τ_n on Aggregate Welfare



To get a sense of how the model's response to a child care subsidy compares with what has been documented from experimental data, I compute the elasticity of child outcomes with respect to the child care

subsidy, using the baseline distribution of families and the baseline price of child care. I find that this ranges from - 0.082 to - 0.129 percent (Figure 8).³⁰ For external validation, I turn to the study of various policy interventions by Chaparro, Sojourner, and Wiswall (2020). This study uses an experimental sample of low-birthweight children, whose families come from a variety of socioeconomic backgrounds, to study the effects of various policy interventions. In Table 5 of that paper, the authors report elasticities of child outcomes at age 3 with respect to various policy interventions. I compare the output of my framework with their value of - 0.08 percent for the elasticity of child outcomes with respect to child care's hourly cost. This is quite close to the range I find, especially at lower levels of the subsidy. I interpret the close alignment of these statistics as further validation of the model's fit to untargeted moments.

Figure 8: External Validation: Comparison with Chaparro et. al. 2020



6.2.2 Counterfactual Effects of a Child Care Subsidy: Randomized Family Type

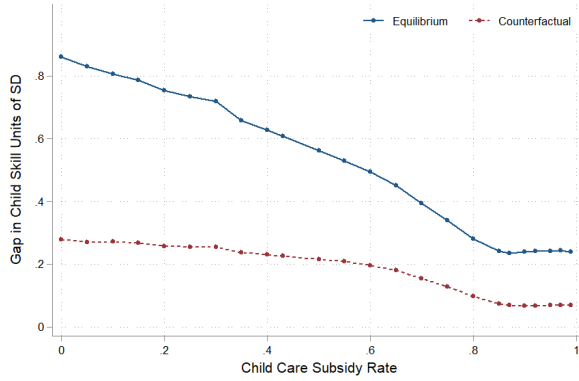
Using the model output, I can construct counterfactual child outcomes in the case where family structure type is exogenously and randomly assigned, rather than being an endogenous choice. This eliminates sorting into family structure by parent skill. The result is that the gap in child skill narrows between single mothers and couples even without a subsidy, because the average income of single mothers increases relative to couples. Performing the same experiment in this environment with random marital status assignment, while holding fixed the total marriage rate at each level of the subsidy, I record the child skill outcomes across family structures. Figure 9a displays counterfactual and equilibrium gaps across family structures in child skill outcomes, over different levels of the subsidy.

In this model, single parenthood is usually the result of a high-skilled man meeting a low-skilled woman and rejecting the match. Women rarely reject men, because single motherhood turns out to be such a difficult state.

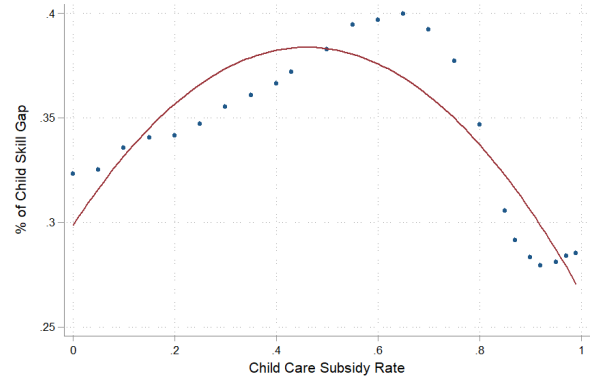
Rejected matches, in equilibrium, would have counterfactually formed a couple with only slightly lower

³⁰This range does not change very much if I also allow the distribution of families and parental skill to adjust in equilibrium.

Figure 9: The Effect of a Child Care Subsidy on Child Outcomes



(a) Countefactual Effect of a Child Care Subsidy τ_n on Child Skill



(b) Technology's Contribution to the Gap in Child Skill Outcomes

income than the average couple which does form. These counterfactual couples would have funded similar investment in children at the expense of lower consumption. On the other hand, single parent households that do not arise in equilibrium would have counterfactually resulted in highly skilled women parenting alone, with much higher income than the one-parent families which arise endogenously. These counterfactual single mothers would have invested more in their children, due to their higher income and parenting productivity. Thus, randomizing the marriage market outcome does not affect the composition of child outcomes among couples, but it does affect the composition of child outcomes among single mothers. Accordingly, controlling for endogenous sorting into family structure reduces the gap in child outcomes across family structures because it raises the outcome of children raised by single mothers. In Figure 9a, this is reflected in the lower gap at any level of the subsidy in the counterfactual economy.

Referring back to the definition of the marginal cost elasticities presented in equations (13) and (15), in equilibrium the income composition of single mothers works to reinforce the elasticity of their technology to the price of child care. Randomly assigning parents to family structures breaks the endogenous link between parenthood skill and single parenthood: the remaining gap is due to exogenous attributes of the two family structures, rather than endogenous composition differences. This contribution from technology, expressed as a fraction of the gap at each level of the subsidy, is presented in Figure 9b.

The investment technology, which incorporates differences in the number of parents, plays a large role in explaining why one-parent families benefit the most from child care subsidies in this model. In particular, technological differences explain between 28 and 40 percent of the gap in child skill, depending on the level of the subsidy.

7 Conclusion

In this paper, I examine the role of family structure in determining the effects of a child care subsidy. In particular, I first specify two investment technologies and empirically test and reject the hypothesis that one-

and two-parent families use the same technology to invest in their child's skill. I find that the investment technology used by single mothers emphasizes non-parental child care more than the one used by couples, although the elasticity of substitution between inputs is statistically the same. Using the elasticity of investment's price with respect to the price of child care, I highlight an implication of the empirical finding for a common policy intervention intended to increase child skill: child care subsidies. The estimation implies that one-parent families will have a larger response to child care subsidies than two-parent families will.

To perform policy experiments, I expand the model to include endogenous family formation and a child care subsidy which is financed by a labor income tax. I document the effects of a child care subsidy in the stationary equilibrium of this expanded model, and verify the implications of the previous estimation: single mothers see higher gains in child skill, driven both by their lower income and by the technology they use to invest in their children. Finally, I use the general equilibrium framework to perform a counterfactual exercise in which endogenous sorting into family structures is shut off. This allows me to decompose the higher gains seen by children of single mothers into a component due to exogenous differences, such as investment technologies, and a residual component due to composition differences in parental skill across the two groups.

References

- Abbott, B. (2020). Incomplete markets and parental investments in children. *Working Paper*.
- Abbott, B., G. Gallipoli, C. Meghir, and G. Violante (2019). Education policy and intergenerational transfers in equilibrium. *Journal of Political Economy* 127(6).
- Agostinelli, F. and M. Wiswall (2017). Identification of dynamic latent factor models: The implications of re-normalization in a model of child development. *NBER Working Paper No. 22441*.
- Aguiar, M. and E. Hurst (2007). Measuring trends in leisure: The allocation of time over five decades. *The Quarterly Journal of Economics* 122(3), 969–1006.
- Almond, D. and J. Currie (2011). Human capital development before age five. In *Handbook of Labor Economics*, Chapter 15, pp. 1315–1486.
- Anderson, M. L. (2008). Multiple inference and gender differences in the effects of early intervention: A reevaluation of the abecedarian, perry preschool, and early training projects. *Journal of the American Statistical Association* 103(484), 1481–1495.
- Andreoni, J. (1990). Impure altruism and donations to public goods: A theory of warm-glow giving. *The Economic Journal* 100(401), 464–477.
- Baker, M., J. Gruber, and K. Milligan (2008). Universal childcare, maternal labor supply, and family well-being. *Journal of Political Economy* 116(4), 709–745.
- Balbo, N., F. C. Billari, and M. Mills (2013). Fertility in advanced societies: A review of research. *European Journal of Population* 29, 1–338.
- Barro, R. J. and G. S. Becker (1989). Fertility choice in a model of economic growth. *Econometrica* 57(2), 481–501.
- Berlinski, S., S. Galiani, and P. Gertler (2009). The effect of pre-primary education on primary school performance. *Journal of Public Economics* 93(1-2), 219–234.
- Blau, D. and J. Currie (2006). Pre-school, day care, and after-school care: Who’s minding the kids? In E. A. Hanushek and F. Welch (Eds.), *Handbook of the Economics of Education*, Chapter 20, pp. 1163–1276. Elsevier.
- Boar, C. (2020). Dynastic precautionary savings. *Working Paper*.
- Boldrin, M., M. De Nardi, and L. E. Jones (2015). Fertility and social security. *Journal of Demographic Economics* 81(3), 261–299.
- Caucutt, E. and L. Lochner (2020). Early and late human capital investments, borrowing constraints, and the family. *Journal of Political Economy* 128(3), 1065–1147.
- Chaparro, J., A. Sojourner, and M. J. Wiswall (2020). Early childhood care and cognitive development. *NBER Working Paper No. 26813*.

- Cunha, F., J. Heckman, and S. Schennach (2010). Estimating the technology of cognitive and noncognitive skill formation. *Econometrica* 78(3), 883–931.
- Daruich, D. (2020). The macroeconomic consequences of early childhood development policies. *Working Paper*.
- De Nardi, M. (2004). Wealth inequality and intergenerational links. *Review of Economic Studies* 71, 743–768.
- Del Boca, D., C. Flinn, and M. Wiswall (2014). Household choice and child development. *Review of Economic Studies* 81, 137–185.
- Department of Education, U. S. (2001). Early childhood longitudinal study, birth cohort (restricted use license).
- Garcia, J. L., J. J. Heckman, and A. L. Ziff (2017). Gender differences in the benefits of an influential early childhood program. *European Economic Review* 109, 9–22.
- Gauthier, A. H. (2007). The impact of family policies on fertility in industrialized countries: A review of the literature. *Population Research and Policy Review* 26(3), 323–346.
- Gayle, G.-L., G. Limor, and M. Soytas (2017). What is the source of the intergenerational correlation in earnings? *Working Paper*.
- Guner, N., R. Kaygusuz, and G. Ventura (2016). Child-related transfers, household labor supply and welfare. *Working Paper*.
- Gupta, N. D. and M. Simonsen (2010). Non-cognitive child outcomes and universal high quality child care. *Journal of Public Economics* 94(1-2), 30–43.
- Güvenen, F. and M. Rendall (2015). Women’s emancipation through education: A macroeconomic analysis. *Review of Economic Dynamics* 18(4), 931–956.
- Havnes, T. and M. Mogstad (2014). Is universal child care leveling the playing field? *Journal of Public Economics* 127, 100–114.
- Heckman, J. J., J. Stixrud, and S. Urzua (2006). The effects of cognitive and noncognitive abilities on labor market outcomes and social behavior. *Journal of Labor Economics*.
- Kottelenberg, M. J. and S. F. Lehrer (2014). The gender effects of universal child care in Canada: Much ado about boys? *Working Paper*.
- Kottelenberg, M. J. and S. F. Lehrer (2017). Targeted or universal coverage? assessing heterogeneity in the effects of universal childcare. *Journal of Labor Economics* 35(3), 609–635.
- Lee, S. Y. and A. Seshadri (2019). On the intergenerational transmission of economic status. *Journal of Political Economy* 127(2), 855–921.

- Lochner, L. J. and A. Monge-Naranjo (2011). The nature of credit constraints and human capital. *The American Economic Review*, 2487–2529.
- McGrattan, E. R. and E. C. Prescott (2017). On financing retirement with an aging population. *Quantitative Economics* 8(1), 75–115.
- NCES. Data collected through the ecls-b: Child assessment data. Found at: https://nces.ed.gov/training/datauser/ECLS-B_03A.html?dest=ECLS-B_03_S0090.html.
- Nellis, L. and B. E. Gridley (1994). Review of the bayley scales of infant development - second edition. *Journal of School Psychology* 32(2), 201–209.
- Nisbett, R. E., J. Aronson, C. Blair, W. Dickens, J. Flynn, D. F. Halpern, and E. Turkheimer (2012). Intelligence: New findings and theoretical developments. *American Psychologist* 67(2), 130–159.
- Restuccia, D. and C. Urrutia (2004). Intergenerational persistence of earnings: The role of early and college education. *American Economic Review* 94(5), 1354–1378.
- Shonkoff, J. P. (2010). Building a new biodevelopmental framework to guide the future of early childhood policy. *Child Development* 81(1), 357–367.

Appendix

Contents (Appendix)

I	Data Source Descriptions	41
A	The Early Childhood Education Longitudinal Study, Birth Cohort	41
A.1	Direct Assessments of Child Skill	41
A.2	Parental Time and Child Care Time	43
A.3	Summary Statistics of the ECLS-B	44
A.3.1	Sumstats Raw ECLS-B Sample	44
A.3.2	Sumstats Raw ECLS-B Sample, Weighted	46
B	Summary Statistics for Imputation Sample in American Time Use Survey	50
II	Empirical Motivation	51
C	Fertility and Direction-of-Transfer Assumptions	52
D	Assumption of One-Shot Marriage Market	54
E	Assumptions on Child Skill and Gender	55
E.1	Correlation of Initial Child Skill and Final Child Skill with Family Income	57
F	Assumptions About Investment Inputs: Child Care Spending vs. Total Spending in the PSID	
	CDS	58
G	Assumptions About Time Inputs From Single Fathers	60

Part I

Data Source Descriptions

A The Early Childhood Education Longitudinal Study, Birth Cohort

The ECLS-B follows a nationally representative sample of families raising a child who was 9 months old in 2001. It was designed and collected by the United States Department of Education. Using birth-certificate data from the National Center for Health Statistics, over 14,000 births were selected within Primary Sampling Units. Children of mothers younger than 15 were excluded from the sampling frame. There are 5 waves: wave 1 is the 9-month old data collection round, wave 2 occurs at 2 years, wave 3 at 4 years, and waves 4 and 5 at kindergarten entry. If the focal child was not in kindergarten when wave 4 was collected, the surveyors went back and collected data the next year when they were enrolled. In addition, if a child repeated kindergarten, their scores were also collected in wave 5 in addition to wave 4. Each wave contains several instruments; these are different self-administered questionnaires (SAQs) for different people in the child's life, in addition to the child-level data. Table 10 summarizes these instruments in each wave of the survey.

Table 10: The Structure of the ECLS-B

Instrument	Wave 1	Wave 2	Wave 3	Wave 4+5
1.	Parent Interview	Parent Interview + SAQ	Parent SAQ	Parent SAQ
2.	Resident Father	Resident Father SAQ	Resident Father SAQ	ECEP Interview ¹
3.	Nonresident Father	Nonresident Father SAQ	Preschool Center Director SAQ	Teacher
4.		Child Care Provider	Preschool ECEP SAQ ¹	WECEP Interview ²
5.		Center Director		

¹ Early Care and Education Provider

² Wrap-around Care Early Care and Education Provider

A.1 Direct Assessments of Child Skill

For each wave of the survey, an age-appropriate assessment of the child's skill is administered (for a detailed overview, see ? and NCES (NCES)). In waves 1 and 2, when the child is 9 months and 2 years old, respectively, the Bayley short form, research edition (BSF-R) cognitive direct assessment was administered. This assessment is constructed using a subset of questions from a standard assessment for children from this age group, the Bayley Scales for Infant Development, 2nd Edition (BSID-II). The BSID-II, which was reviewed in detail by Nellis and Gridley (1994), is considered an excellent tool for assessing children 42 months of age and younger. The BSF-R in turn is designed to collect the same information as the BSID-II, but with fewer questions, making it easier to administer and score. This was necessary because, unlike the BSID-II, the BSF-R was administered by interviewers who were not experts in child psychology for the purpose the ECLS-B. Next, the results of the BSF-R were mapped into the metric of the BSID-II. There

are mental and motor scores assigned to children who completed the test in each wave. These scores have different ranges: for mental, the range is 0-178, and for motor the range is 0-111. The BSF-R is organized as follows. The child is administered a series of questions, where their performance determines the total set of questions they are ultimately asked. The questions begin with a basal set, where the questions are grouped by the skill they are meant to measure, where the skills are those children in the age group are expected to demonstrate. If the child performs poorly, they continue to a basal set of questions. If they perform well, they are administered a ceiling set of questions. The mental knowledge and skill of the child are reported in the mental score, which is reported as a scale score, a t-score, and a probability score. The scale score can be compared across waves or within a cross-section, but the t-score is better suited to comparing across groups within a cross-section. The probability score measure the probability the child has acquired the skills measured by the test. For the correlations of child skill with family attributes, and the regressions using child skill, I use the mental scale score.

The skills measured by the mental scale score are arranged in ascending order of development:

1. Explores objects (i.e. reaching for and holding objects, manipulating them, and banging them)
2. Explores purposefully
3. Jabbers expressively
4. Early problem solving
5. Names objects
6. Receptive vocabulary
7. Expressive vocabulary
8. Listening/comprehension
9. Matching/discrimination

Because there was existing single measure of skill that be used continually to assess child skill during the first five years of life, a change to a new assessment was necessary for the preschool and kindergarten collections. In addition, no existing assessment was age-appropriate for the third and fourth wave of the ECLS-B while also being straightforward to administer. Therefore, a new cognitive assessment was developed especially for the ECLS-B, the cognitive battery. The content of this battery included a reading assessment and a mathematics assessment. For the preschool (wave 3) round there was a color knowledge assessment as well, to test for children's knowledge of basic colors. The reading assessment "Examines children's letter recognition, letter sound knowledge, recognition of simple words, phonological awareness, receptive and expressive vocabulary knowledge, and knowledge of print conventions", while the mathematics assessment "Examines children's number sense, counting, operations (e.g., addition, subtraction, multiplication, division), geometry, pattern understanding, and measurement".

These assessments were administered in a way structured to the child's ability, so that children who did

not demonstrate sufficient english language skill were not administered the literacy portion (for example). That means that not all of the questions in the assessment were asked of each child. The types of scores are provided for the reading assessment: theta scores and overall scale scores. The theta score is normal distributed and ranges from -2 to 2; it estimates the child's score if they had been administered all of the assessment. The scale score estimates the number of items the child would have gotten correct on the entire assessment. Both score are appropriate for analysis of children's rank within a cross-section and changes across time. The mathematics assesment is structured similarly to the reading assessment and also reports the theta score and the scale score. I used the scale score for the reading and mathematics assessment and take the average across the two assessments to construct a single score for the child.

A.2 Parental Time and Child Care Time

In each wave of the survey, the primary care provider (usually the mother) and the resident father fill out detailed questionnaires on the activities they do with their kids and at what frequency (once a week, twice a week, once a month, etc.). In addition, they report age, educational attainment, income, hours worked, the number of hours the child spent in non-parental care, what type of care that was (relative, non-relative, center-based), and the cost of that care.

I define "quality time" as the total amount of time spent (1) reading to the child, or (2) playing outside with the child. To map from frequencies of activities to levels of quality time supplied by parents, I impute amount of time per activity using data from the ATUS. The imputation uses common characteristics observed across both samples: gender, marital status (married/cohabiting or single), labor force status, and educational attainment. Here educational attainment is defined as being either having less than a college degree, or having a college degree or more. For hourly wages, I use time spent working and income to compute the pre-tax levels, and then Table 2 of McGrattan and Prescott (2017) to correct for labor income taxes. For hourly prices of non-parental care, I use total cost of child care and total hours in child care for the primary source of non-parental care reported by the primary caregiver of the survey child subject.

In the following section I report sample summary statistics for the raw ECLS-B sample, before I impose restrictions on it for the estimation sample. The population moments I use in the internal calibration for the fraction of parents who are single mothers comes from this sample. The fraction in the sample that also reports variables necessary for estimation is larger than these population moments. Notice that the fraction below 185% of the poverty line in the pooled sample is quite high, at 50%. This drop to 40% by the time the child is age 4. Averages here include observations for which the response is 0. This explains why the average age of the father is now lower than the mother's.

A.3 Summary Statistics of the ECLS-B

A.3.1 Sumstats Raw ECLS-B Sample

Table 11: ECLS-B Data Moments (Waves 1-3, Unweighted)

	count	mean	Levels			sd	min	max
			p10	p50	p90			
Hours Ed. Time: Mother	32050	5.83	0.00	3.84	18.91	6.51	0.00	22.76
Hours CC (Primary)	32050	15.00	0.00	6.00	40.00	17.99	0.00	120.00
Hours CC (Combined)	31200	17.07	0.00	8.00	45.00	19.87	0.00	144.00
Total Cost CC (Primary)	32050	1.12	0.00	0.00	3.75	2.85	0.00	99.50
Total Cost CC (Combined)	17550	2.08	0.00	1.11	5.00	3.58	0.00	99.50
After-Tax Wages: Mother	28250	11.21	4.60	8.89	18.55	17.65	0.00	1063.80
$\frac{p_n}{w_m}$	28250	76.95	0.00	0.00	0.35	8703.86	0.00	1345064.75
$\frac{HoursCC(Primary)}{HoursEd.Time:Mothers}$	17100	0.45	0.02	0.16	1.08	0.99	0.00	22.76
Family Income	29450	56407	10000	37500	150000	558423	1	300000
Family Income: After Tax	29450	53753	20359	41413	117460	38769	11763	223092
Age: Primary Caregiver W1	29350	30.21	21.00	30.00	39.00	6.92	15.00	82.00
Age: First Child (Mother)	28400	23.85	17.00	23.00	32.00	6.04	10.00	50.00
Rates								
B.A. Mother	32050	0.33						
Below 100% Poverty Line	32050	0.23						
Below 185% Poverty Line	32050	0.44						
Single Mother	32050	0.17						
Obs_rounded	32050							

Source: U.S. Department of Education, National Center for Education Statistics,
Early Childhood Longitudinal Study, Birth Cohort (ECLS-B),
Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File.
All sample counts have been rounded to the nearest 50 in accordance with NCES requirements.

In Table 11, the unweighted summary statistics for a pooled sample of waves 1-3 are presented. Over these three waves, mothers on average spend 6 hours of their time per week investing in their children. The mean of this variable is higher than the median, which is about 3.8 hours per week. The mean is more sensitive to outliers—in this case, the high-value outliers are driving the higher value for the mean. Most families don't spend large numbers of hours per week investing parental time in their children. Meanwhile, the total hours spent in the primary source of child care are on average 15, which is much higher than the median of 6 hours per week. This reflects the fact that in the raw sample some families report that their child spend 120 hours per week in child care. Such observations are of interest because they represent the other corner of the investment decision: not zero investment, but rather investment at a binding time constraint of the child. This table demonstrates that very few families are at this constraint in the equilibrium of the United States in 2001, which generated the data documented in the ECLS-B.

If, instead of restricting to the primary source of child care, I looked at the hours spent in any form of non-parental care, the average moves up to 17 and the median to 8. The maximum value for this category is 144 hours per week, which requires spending more than 20 hours a day in non-parental child care. In the estimation sample I only use observations for families that use less than 100 hours total of their child's time in parental investment activities or non-parental child care. This is because the estimation equations I use are only valid if the time constraint of the child is slack.

Moving on to the next category of variables, the price per unit of the time inputs into investment is documented in rows 5-7 of this table. The price for the primary source of child care is lower than for all child care combined. The latter variable is the average across all sources of child care - it is higher because some families use small amounts of expensive child care. So, even if the hours in child care don't increase much moving from primary to total sources of child care, the price does jump. The median price per hour of child care is 0. This is because many families use no child care and many also use child care that is completely free. My estimation assumptions effectively assume that child care which is free is not an input into investment.

The ratio of the price of child care to the mother's wage is high on average. The moment reported here uses observations with very small values of the denominator, which are excluded in the estimation sample. Moving on to income, before-tax income is on average \$56,000.³¹ Once I implement the method I use to correct for taxes, using McGrattan and Prescott (2017), the redistributive nature of the tax system that I assume is evident from the fact that, while the mean decreases, the median increases and so does the minimum income in the sample. When I report the correlation of child skill and family income over time during the calibration section of the model parameterization section in the main body of this paper, I am using after-tax income.

Finally, the last section of the levels statistics for the ECLS-B shows statistics on the age of the primary caregiver at the time of the interview, and on the mother's age when her first child was born. If the primary caregiver is the biological mother, as I impose in my estimation sample, then comparing these two rows would give some intuition for whether the child in the ECLS-B sample is the mother's first child. Here, however, they can be different people.

The second part of Table 11 deals with rates in the raw sample. The rate of BA attainment for the primary caregiver is about 1 in 3. The poverty rate is about 1 in 4, and the unweighted rate of single motherhood is about 17%.

³¹ This is between the mean and median family income for families with children under 18 in current dollars for the years 2001-2004, as reported by the United States Census Bureau in Table F-9 here: [census.gov/data/tables/time-series/demo/income-poverty/historical-income-families.html](https://www.census.gov/data/tables/time-series/demo/income-poverty/historical-income-families.html)

A.3.2 Sumstats Raw ECLS-B Sample, Weighted

Table 12: ECLS-B Data Moments (Waves 1-3, Weighted)

	count	mean	Levels			sd	min	max
			p10	p50	p90			
Hours Ed. Time: Mother	29450	6.52	1.10	3.84	20.01	6.58	0.00	22.76
Hours CC (Primary)	29450	16.04	0.00	9.00	40.00	17.77	0.00	120.00
Hours CC (Combined)	29450	17.71	0.00	10.00	45.00	19.50	0.00	144.00
Total Cost CC (Primary)	29450	1.27	0.00	0.00	4.00	2.92	0.00	99.50
Total Cost CC (Combined)	17500	2.16	0.00	1.25	5.14	3.51	0.00	99.50
After-Tax Wages: Mother	28200	10.92	4.56	8.60	18.21	16.53	0.00	1063.80
$\frac{p_n}{w_m} \text{HoursCC(Primary)}$	28200	144.23	0.00	0.00	0.37	12790.67	0.00	1345064.75
$\frac{\text{HoursEd.Time:Mothers}}{\text{HoursCC(Primary)}}$	17050	0.47	0.03	0.17	1.13	0.96	0.00	22.76
Family Income	29450	55624	10500	37500	150000	54483	1	300000
Family Income: After Tax	29450	53233	20736	41413	117460	37820	11763	223092
Age: Primary Caregiver W1	29300	30.01	22.00	30.00	39.00	6.72	15.00	82.00
Age: First Child (Mother)	26200	23.73	17.00	23.00	32.00	5.82	10.00	50.00
Rates								
B.A. Mother	29450	0.25						
Below 100% Poverty Line	29450	0.24						
Below 185% Poverty Line	29450	0.47						
Single Mother	29450	0.19						
Obs_rounded	29450							

Source: U.S. Department of Education, National Center for Education Statistics, Early Childhood Longitudinal Study, Birth Cohort (ECLS-B), Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File.
All sample counts have been rounded to the nearest 50 in accordance with NCES requirements.

In Table 12, the weighted summary statistics for a pooled sample of waves 1-3 are presented. I use cross-section survey weights for each observation and report pooled moments in this table. Most of the statistics are very similar after weighting; only slight changes to means are noticeable. One exception is the rate of single motherhood, which is 19% in the weighted sample (versus 17% in the unweighted sample). In the model calibration section, I use the marriage rate of the weighted total sample as a target, although I use the time investment levels of the estimation sample for targets as well.

Effect of sample cleaning on composition of child care type used in the sample

In the process of sample cleaning, the composition of child care sources shifts. In Table 13 I report the raw sample distribution for primary source of child care over child care type. By comparing this table with Table 16 and Table 15, I can describe how the estimation sample differs from the population in terms of the type of child care the family uses as their primary source. In particular, when children are very young many families report not using any non-parental care (wave 1), and in addition some sources of child care often do

cost anything (like care from relatives). These points are illustrated by comparing Table 15 with Table 13.

Table 13: Primary Source of Child Care by Wave: Raw Sample

	(1)	(2)	(3)
	Wave 1	Wave 2	Wave 3
	pct	pct	pct
NOT ASCERTAINED	0.19	0.13	0.22
NO NONPARENTAL CARE	50.04	50.55	18.54
RELATIVE CARE IN CHILDS HOME	12.61	8.60	5.39
RELATIVE CARE IN ANOTHER HOME	12.52	9.96	6.39
RELATIVE CARE, LOCATION VARIES	1.51	0.95	0.77
NONRELATIVE CARE IN CHILDS HOME	4.00	3.11	1.64
NONRELATIVE CARE IN ANOTHER HOME	10.38	11.02	5.47
NONRELATIVE CARE, LOCATION VARIES	0.12	0.23	0.18
CENTER-BASED PROGRAM	7.86	14.95	46.34
EQUAL TIME IN MULTIPLE ARRANGEMENTS	0.79	0.49	1.99
HEAD START PROGRAM			13.06
Total	100.00	100.00	100.00
Obs_rounded	10700	9850	8950

Source: U.S. Department of Education, National Center for Education Statistics,
Early Childhood Longitudinal Study, Birth Cohort (ECLS-B),
Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File.

All sample counts have been rounded to the nearest 50 in accordance with NCES requirements.

In Table 14, I report similar moments from the ECLS-B on type of child care used in the raw sample, by family structure. For each family structure, the observations are pooled across waves 1 through 3. In this table, it is evident that single mothers use more non-parental relative care than couples do. Since relative care may be paid for less often, one concern with the estimation approach in this paper is that I am dropping information on an input into investment in a way which disproportionately affects single mothers, because they are more likely to use free child care (for example, care from relatives). In the discussion based on this Table, I address the general concern about free child care services and the sample restrictions I impose on the ECLS-B data with respect to those observations. Recall that in the estimation sample, I exclude observations that pay less than 1 cent per hour for their primary source of child care.

Consider the estimation result that one- and two-parent families use different technologies. This result is driven by the relative quantities of mother time and child care time being very different across family structures, despite facing similar relative prices. It is possible that free child care is used by single mothers in a way that, if I could observe relative “prices” for these observations, would make single mothers look more like couples on average. Then including this source, if we could observe its “price”, would change the result that the technologies are statistically different.

In order for including these observations to change the main result, families whose primary source of non-parental child care is free would have to select a ratio of quantities very different from the estimation sample

given a similar ratio of prices, or choose a ratio of quantities very similar from the estimation sample for a very different ratio of prices. Here, think of the ratio of prices as using the marginal utility of leisure of the child care provider in place of a price, as described in Blau and Currie (2006).

I can check that the median ratio of quantities is similar for mothers who use free child care: for single mothers in this group, the median ratio of quantities is 0.08, and for cohabiting mothers it is 0.26. Compare this with the estimation sample counterpart medians of 0.07 and 0.23, respectively: they are quite close.³²

If I could observe the relative prices of inputs for families that use free child care, these observations would bring the parameters estimated for one- and two-parent families closer together if the ratio of prices was much lower for one-parent families than for two-parent families. Consider the denominator of this ratio, the mother's wage, which I do observe for both one- and two-parent families using free child care. The median wage for single mothers using free child care is 8 dollars per hour, and for cohabiting mothers it is 16. This is very similar to the estimation sample values of 7.9 and 16.1, respectively. With a lower denominator for the price ratio, a lower ratio for one-parent families would require that the marginal utility of leisure for the provider of free child care (the numerator) be much lower in the one-parent group than the two-parent group. This seems unlikely: for concave utility, it would require that providers of free child care for one-parent families enjoy higher levels of leisure than those of two-parent families. One-parent families demand more child care than two-parent families (33 hours per week vs. 26) in the group using free child care, and if this free care is provided by a relative, that relative is probably more likely to be poor. They are unlikely to be able to finance much higher levels of leisure.

³²I compare medians because relaxing the cutoffs on input prices introduces outliers, which misleadingly affect the mean. The intuition based on means in the estimation sample, outlined in the introduction and estimation section, is also true if one compares medians across family structures in the estimation sample.

Table 14: Primary Source of Child Care by Family Type: Raw Sample (Weighted)

	(1)	(2)	(3)
	Single Mothers	MC Mothers	MC Fathers
	pct	pct	pct
NOT ASCERTAINED	0.23	0.13	0.13
NO NONPARENTAL CARE	30.70	43.87	43.87
RELATIVE CARE IN CHILDS HOME	15.65	7.44	7.44
RELATIVE CARE IN ANOTHER HOME	12.29	9.25	9.25
RELATIVE CARE, LOCATION VARIES	1.08	1.07	1.07
NONRELATIVE CARE IN CHILDS HOME	1.51	3.46	3.46
NONRELATIVE CARE IN ANOTHER HOME	8.97	9.18	9.18
NONRELATIVE CARE, LOCATION VARIES	0.25	0.15	0.15
CENTER-BASED PROGRAM	21.40	21.72	21.72
EQUAL TIME IN MULTIPLE ARRANGEMENTS	1.57	0.90	0.90
HEAD START PROGRAM	6.34	2.82	2.82
Total	100.00	100.00	100.00

Source: U.S. Department of Education, National Center for Education Statistics,
Early Childhood Longitudinal Study, Birth Cohort (ECLS-B),
Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File.

In Table 15, only families reporting a source of non-parental child care who also report paying at least 1 cent per hour for it (and several other criteria imposed on the estimation sample) are shown. Center based programs are especially overrepresented relative to the raw sample, especially in the first wave. So is non-relative care in another home. Care from relatives, meanwhile, is used by a smaller percentage of the population in the estimation sample, relative to the raw sample.

Table 15: Primary Source of Child Care: Estimation Sample

	(1)	(2)	(3)
	Wave 1	Wave 2	Wave 3
	pct	pct	pct
RELATIVE CARE IN CHILDS HOME	4.10	5.35	2.51
RELATIVE CARE IN ANOTHER HOME	11.05	7.74	4.33
RELATIVE CARE, LOCATION VARIES	0.34	0.23	0.11
NONRELATIVE CARE IN CHILDS HOME	11.85	9.34	5.58
NONRELATIVE CARE IN ANOTHER HOME	43.17	39.52	21.64
NONRELATIVE CARE, LOCATION VARIES	0.68		0.34
CENTER-BASED PROGRAM	28.82	37.81	63.67
HEAD START PROGRAM			1.82
Total	100.00	100.00	100.00
Obs_rounded	900	900	900

Source: U.S. Department of Education, National Center for Education Statistics,
Early Childhood Longitudinal Study, Birth Cohort (ECLS-B),
Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File.

In Table 16, the use of different primary child care sources over family types is tabulated for the estimation sample. Comparing these distributions across family types within the estimation sample, single mothers differ from other family type primarily in lower use of nonrelative care in child's home and in their higher use of head start programs. The latter point is not surprising: head start is means tested, and single mothers in the United States are overrepresented among the poor. Another point that is apparent is that, although I do not require the married/cohabiting mothers to have a resident father who completes a resident father survey, the distribution for the two groups (resident mothers and resident fathers) is nevertheless very similar.

Table 16: Primary Source of Child Care: Estimation Sample

	(1)	(2)	(3)
	Single Mothers	MC Mothers	MC Fathers
	pct	pct	pct
RELATIVE CARE IN CHILDS HOME	6.45	3.75	3.90
RELATIVE CARE IN ANOTHER HOME	7.53	8.06	7.01
RELATIVE CARE, LOCATION VARIES	1.08	0.14	0.20
NONRELATIVE CARE IN CHILDS HOME	1.61	9.38	9.61
NONRELATIVE CARE IN ANOTHER HOME	33.87	34.51	35.54
NONRELATIVE CARE, LOCATION VARIES		0.28	0.50
CENTER-BASED PROGRAM	46.77	43.54	42.64
HEAD START PROGRAM	2.69	0.35	0.60
Total	100.00	100.00	100.00
Obs_rounded	200	1450	1000

Source: U.S. Department of Education, National Center for Education Statistics,
Early Childhood Longitudinal Study, Birth Cohort (ECLS-B),
Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File.

B Summary Statistics for Imputation Sample in American Time Use Survey

The time per activity for each demographic bin, as computed in the ATUS, is reported in Tables 17 and 18. The number of observations per bins is also reported in Table 19.

Table 17: ATUS Imputation Sample (Unweighted)

	count	mean	p10	p50	p90	sd	min	max
Edited: sex	8395	1.69	1.00	2.00	2.00	0.46	1.00	2.00
Edited: age	8395	34.26	27.00	34.00	41.00	5.76	15.00	54.00
Number of own children < 18 years of age	8395	2.12	1.00	2.00	3.00	1.03	1.00	10.00
time_cc_play_hrs	4616	1.82	0.50	1.50	3.92	1.44	0.08	10.30
time_cc_reading_hrs	6729	0.49	0.17	0.42	1.00	0.36	0.03	4.50
indicator_gtHS	8395	0.83	0.00	1.00	1.00	0.38	0.00	1.00
indicator_gteBA	8395	0.60	0.00	1.00	1.00	0.49	0.00	1.00
indicator_couple	8395	0.93	1.00	1.00	1.00	0.25	0.00	1.00
income	3562	123795.94	22500.00	87500.00	383150.00	121383.93	2500.00	383150.00
Observations	8395							

These moments are from the pooled 2003-2016 ATUS sample. Moments are weighted with ATUS final weights.

For parents aged 15-55, living with own child aged 3 or less

Table 18: ATUS Imputation Sample (Weighted)

	count	mean	p10	p50	p90	sd	min	max
Edited: sex	8395	1.67	1.00	2.00	2.00	0.47	1.00	2.00
time_cc_reading_hrs	6729	0.48	0.17	0.42	1.00	0.35	0.03	4.50
time_cc_reading_hrs	6729	0.48	0.17	0.42	1.00	0.35	0.03	4.50
time_cc_play_hrs	4616	1.72	0.50	1.25	3.75	1.45	0.08	10.30
indicator_gtHS	8395	0.76	0.00	1.00	1.00	0.42	0.00	1.00
indicator_gteBA	8395	0.53	0.00	1.00	1.00	0.50	0.00	1.00
indicator_couple	8395	0.92	1.00	1.00	1.00	0.27	0.00	1.00
income	3562	120722.31	22500.00	87500.00	383150.00	121987.69	2500.00	383150.00
Observations	8395							

These moments are from the pooled 2003-2016 ATUS sample. Moments are weighted with ATUS final weights.

For parents aged 15-55, living with own child aged 3 or less

Table 19: ATUS Time per Activity Averages by Demographic Bins

Bin	Couple	Gender	gt HS	Reading Hrs	Ed Hrs	Talk. Hrs	Play Hrs	CC Hrs	N Reading	N Ed	N Talk.	N Play	N CC
1	0	1	0	0.793	0.801	0.745	4.222	4.115	7	21	15	11	21
2	0	1	0	.	0.477	0.477	.	3.229	0	4	4	0	4
3	0	1	1	0.519	0.434	0.271	0.566	1.600	12	17	7	4	17
4	0	1	1	0.500	0.500	.	4	5	1	1	0	1	1
5	0	2	0	0.746	0.745	0.672	1.322	2.926	85	147	74	62	147
6	0	2	0	0.387	0.453	0.491	2.353	3.300	39	75	40	37	75
7	0	2	1	0.416	0.592	0.780	1.287	2.766	185	245	76	109	245
8	0	2	1	0.664	0.721	0.651	1.766	4.303	45	69	34	36	69
9	1	1	0	0.503	0.560	0.557	1.622	2.104	232	367	157	133	367
10	1	1	0	0.310	0.471	0.690	1.617	5.031	10	18	8	8	18
11	1	1	1	0.454	0.474	0.415	1.508	2.256	1745	2101	482	1117	2101
12	1	1	1	0.585	0.656	0.685	1.574	3.638	71	77	14	63	77
13	1	2	0	0.474	0.569	0.596	1.562	2.904	239	390	182	153	390
14	1	2	0	0.494	0.608	0.582	1.743	4.210	302	447	212	269	447
15	1	2	1	0.451	0.528	0.550	1.680	3.324	2397	2814	704	1597	2814
16	1	2	1	0.549	0.634	0.533	2.101	4.554	1359	1602	500	1016	1602

Part II

Empirical Motivation

C Fertility and Direction-of-Transfer Assumptions

The way fertility is defined in this model can be separated into two categories: the extensive margin and the intensive margin. I assume that all families have children (the extensive margin), and I assume that families have 2 children each (the intensive margin). In addition, I assume that parents are altruistic and make transfers to children; I do not allow for children to be altruistic and make transfers to their parents. In this section, I offer supporting evidence for each of these three assumptions.

For the first point, I focus on the general inconclusiveness of the literature on the subject. In addition, even for studies which find an effect, it is important to keep in mind the low responsiveness measured along this margin in response to various policies (not just child care subsidies).

For the second point, I provide summary statistics from the ECLS-B which show similar average number of children across one- and two-parent families. For both family structures, the average is about 2 children each.

For the third point, I discuss the competing motivations for having children in the economics literature, to put my decision to model altruistic parents into a broader context.

Assumption that Fertility is Exogenous

There is a sizeable literature on how policy affects fertility decisions. This literature has not reached a definitive conclusion: results tend to be country-specific, and for those that do find an effect, the magnitude is small. One important margin of response to these policies seems to be the timing of birth over the mother's life, rather than the total number of births (total lifetime fertility). For surveys of the literature on this topic, see Gauthier (2007).³³

Obviously, the fact that fertility is exogenous in the model developed within this paper makes it unresponsive to child care subsidies. If, instead, fertility had a strong response to the policies analyzed here, it is possible that the results could be dampened.³⁴ Underlying concerns about this source of invalidation is the classic quality-quantity tradeoff that underlies the decision to have more children in many economic models of fertility.

In my model, quality of children (their skill level) is the margin by which families respond to the child care subsidy, rather than quantity. I invoke a lack of clear evidence for quantity responsiveness to these policies as part of my motivation for modelling it in this way. By contrast, there is sizeable evidence documenting a

³³ An interdisciplinary review of studies on the determinants of fertility is provided in Balbo, Billari, and Mills (2013).

³⁴ Given the magnitudes of fertility elasticities found in empirical analyses reviewed in Gauthier (2007), it is unlikely that the sign of the effect of child care subsidies on child skill or welfare would change sign.

responsiveness of child skill (quality) to child care subsidies —especially for poor and one-parent families, the majority of evidence points to a positive response (see the citations in the introduction for relevant studies on this last point).

Assumption of 2 Children Per Family

I make the assumption about number of children in order to achieve replacement rates for the population. I claim that this is not strongly counterfactual: in fact, it is broadly in line with the moments I compute in the ECLS-B.

In particular, in the ECLS-B one- and two-parent families have similar, but not identical, distributions of children living with them at the time of data collection. In Table 2, I report statistics on the distribution for number of children for single mothers and married or cohabiting couples. The means are very close: 2 for single mothers, and 1.8 for couples. The medians of these distributions, however, are not close: the median number of children being raised by couples in the ECLS-B is 1, and for single mothers it is 2.

TABLE HERE

I could also achieve a stable population across generations by allocating the same total number of children asymmetrically across family structures, for example. I do not consider the heterogeneity in number of children that I see in the data sufficient to justify this, but it would be interesting to see how such an extension affects the results. Given my assumption about non-rivalness of family investments, and the fact that single mothers tend to have slightly more children in the data, I anticipate that the welfare gains from a child care subsidy would be even higher, because more children would be raised by single mothers and thus benefit the most from such a policy.

Assumption that Parents Make Altruistic Transfers to Children

Related to model assumptions about fertility, assume that fertility is incorporated somehow into one's framework. Even then, whether quality or quantity of children responds to a child care subsidy is likely to depend on how children affect one's value function: are parents altruistic, so that better child outcomes make them happier, or are children altruistic, so that better child outcomes result in higher transfers to parents (thus making parents happier).

One example of a policy which does appear to have a strong empirical relationship with fertility is social security benefits, as documented and analyzed by Boldrin, De Nardi, and Jones (2015). In that study, the authors evaluate the two competing models of how children affect one's value function described above: one with altruistic parents who make transfers to children (Barro and Becker (1989)), and one with altruistic children who make transfers to parents once the children are adults. Only the latter specification is able to match the negative correlation in the data between social security programs and fertility. In the latter model, social security replaces children, and fertility declines accordingly.

I assume that parents make transfers to children: that is, in my specification investments in children are

motivated by parental altruism. This is the assumption in many other papers in the macroeconomics literature analyzing early childhood investments made by families, including Lee and Seshadri (2019) and Daruich (2020). Boldrin, De Nardi, and Jones (2015) do not try to examine early childhood investments, so they do not evaluate whether future transfers from altruistic children to parents are sufficient to explain the investments in children which we see in the data. For an example of a framework in which both parents and children make transfers to each other, see Boar (2020). In that study, the two-way transfer environment results in a strategic interaction between generations in the same dynasty which is not trivial to solve.

D Assumption of One-Shot Marriage Market

In the general equilibrium framework presented in the main text of this paper, families are formed on a one-shot marriage market. In addition, in the estimation of the skill investment technologies I impose that valid observations must have the same family structure for the first three waves of the survey. I claim that both the modelling assumption and the estimation restriction are good approximations of the data and represent the majority of families parenting young children in the ECLS-B. In Table 20, I present transition matrices for family structure across waves of the ECLS-B. This table shows that family structure is extremely persistent, especially marriage, when children are very young.

In Panel A, I report the probability that a family which is married, never married, or previously married in wave 1 is in any of those statuses in wave 3. The states of being married or never married are extremely persistent: 93% of married couples remain married until their child is 4 years old, and 80% of single mothers remain single mothers for the same period. A more volatile category is being previously married, which means that you are divorced or separated in wave 1. For that category, 31% have become married by the time their child is 4 years old. Despite this, for each of the marital status categories, the portion that persists in that category through wave 3 is by far the majority.

In Panel B, I report the same statistics as Panel A, except that they refer to the transition probabilities from wave 1 to wave 2. Panel C reports these statistics for the transition from wave 2 to wave 3. In both Panel B and Panel C, the pattern is very similar to Panel A. There is by far more churn for those previously married, and very high persistence for other marital statuses.

Table 20: Mother's Marital Status Over Time in the ECLS-B

Panel A: Wave 3 (rows) vs. Wave 1 (columns)

	Married	Never Married	Previously Married
Married	0.93	0.17	0.31
Never Married	0	0.80	0
Previously Married	0.07	0.03	0.69

Panel B: Wave 2 (rows) vs. Wave 1 (columns)

	Married	Never Married	Previously Married
Married	0.98	0.07	0.15
Never Married	0	0.92	0.01
Previously Married	0.02	0.01	0.84

Panel C: Wave 3 (rows) vs. Wave 2 (columns)

	Married	Never Married	Previously Married
Married	0.95	0.11	0.18
Never Married	0	0.87	0
Previously Married	0.05	0.01	0.82

Source: U.S. Department of Education, National Center for Education Statistics, Early Childhood Longitudinal Study, Birth Cohort (ECLS-B), Longitudinal 9- Month-Kindergarten 2007 Restricted-Use Data File.

Table 20 tells a consistent story: for the duration of early childhood (before age 5) family structure is extremely persistent.

E Assumptions on Child Skill and Gender

In Tables 21 and 22, I report regression analyses I use to motivate two modelling assumptions: parents do not target investments by child gender, and initial skill endowments (at 9 months) affect skill outcomes later in life (at 4 years of age).

Investment is not targeted by child gender

Table 21 reports four models, each with a time input choice as the dependent variable. The first two are for married couples, the second two for single mothers. Time investments are predicted by attributes of the parents (hourly wages and educational attainment) and attributes of the child (current skill). Child gender is not a statistically significant predictor of parental time inputs, according to Table 21 . There is some evidence in other studies that parenting behavior and treatment effects of the program vary by the gender of the child (see Garcia, Heckman, and Ziff (2017), Kottelenberg and Lehrer (2014)), but I do not see parenting investment decisions depending on gender in my empirical analysis.

Table 21: Time Investments by Child Gender

	Married Couples		Single Mothers	
	(1)	(2)	(3)	(4)
	Tot. Parental Time	N Time	N Time	Total Time
Child is Female	-0.0845 (0.212)	-0.216 (0.782)	-1.156 (0.850)	0.0434 (0.499)
Child Test Score [0,1]	8.528*** (0.598)	8.281*** (2.219)	1.023 (2.477)	9.945*** (1.389)
B.A.: Father	2.104*** (0.250)	-0.895 (0.904)		
B.A.: Mother	1.850*** (0.247)	2.745** (0.881)	1.578 (1.473)	3.484*** (0.556)
Hourly Wage: Father	0.0105 (0.00588)	-0.0248* (0.0118)		
Hourly Wage: Mother	0.00922* (0.00369)	-0.0146 (0.0131)	-0.0638** (0.0244)	-0.0122 (0.0107)
Constant	-2.921*** (0.824)	10.59*** (3.134)	29.68*** (3.443)	18.85*** (1.905)
R^2	.13	1.1e-02	5.0e-03	1.7e-02
Obs.				

Standard errors in parentheses. N stands for non-parental child care.

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Sample sizes rounded to nearest 50, following NCES requirements.

Measures of initial skill are informative

Table 22 reports three models: for married couples, for single mothers, and for the pooled sample. The dependent variable in all three models is the final skill of the child at age 4. Explanatory variables include the initial skill of the child, gender of the child, indicators for parental educational attainment (BA or higher), and parental hourly wages. Initial test scores are statistically significant predictors for final test scores, and so are parental attributes related to their skill. This motivates including heterogeneity in initial skill endowments in the model, and supports the assumption concept of child skill evolving endogenously in response to parental decisions, which in turn depend on their individual attributes (measured indirectly here with wages and education).

Table 22: Predicting Final Skill with Initial Skill + Parental Attributes

	(1) Married Couples	(2) Single Mothers	(3) All
Initial Test Score (9 Mo.): Stdzd	0.138*** (0.0389)	0.139*** (0.0350)	0.130*** (0.0373)
Hourly Wage: Mother	0.0565 (0.0359)	0.217*** (0.0500)	0.123** (0.0383)
Hourly Wage: Father	0.0971** (0.0312)		
Child is Female	0.108 (0.0619)	0.135 (0.0738)	0.0941 (0.0636)
B.A.: Mother	0.291*** (0.0744)	0.662*** (0.113)	0.546*** (0.0651)
B.A.: Father	0.441*** (0.0788)		
Constant	0.851*** (0.183)	0.654*** (0.160)	1.005*** (0.180)
R^2	0.1695	0.1273	0.1237
Observations	2900	1400	2900

Initial skill has predictive power. Units: standard deviations, except for indicators

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Sample sizes rounded to nearest 50, following NCES requirements.

E.1 Correlation of Initial Child Skill and Final Child Skill with Family Income

Table 23 reports a slightly different version of the same qualitative points made with Tables 21 and 22. The dependent variables in the two models are initial skill and final skill, with both family structures pooled. Initial income at 9 months has no predictive power for the initial skill score. At age 4, however, final skill can be predicted with income (both at age 9 months and 4 years) and the initial test score. The calibration target moments used in the expanded model framework use correlation coefficients, but the regression reported here makes the same qualitative point.

Table 23: Correlations of Skill and Family Income

	(1) Test Score W1	(2) Test Score W3
Family Income W1	0.000114 (0.000103)	0.000651* (0.000254)
Flag: Present in model 2 sample	-0.00677 (0.00754)	
Family Income W3		0.000693*** (0.000204)
Test Score W1 (SD)		0.141** (0.0482)
Constant	1.453*** (0.00733)	1.004*** (0.0705)
R^2	.003	.125
Observations	1300	1500

Income in thousands of dollars. Test scores in standard deviation units.

Standard errors in parentheses.

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Sample sizes rounded to nearest 50, following NCES requirements.

F Assumptions About Investment Inputs: Child Care Spending vs. Total Spending in the PSID CDS

In the literature on estimating skill accumulation technologies during early childhood, it is common practice to include money spent on goods as one of the components of investment (examples include Lee and Seshadri (2019), Daruich (2020) and Abbott (2020)). By contrast, my specification includes time spent in child care instead of money spent on the child. In this section, I use tabulations from the 2001 PSID and 2002 PSID CDS to show how child care expenses contribute to total expenditures on the child. To do this, I construct four different measures of total expenditures on the child (Definitions 1 to 4 in the tables below, with each definition specified in the table footnote). Next, I find the fraction of each measure of total expenditures that comes from spending on child care. I report these fractions in Tables 24 to 27. My conclusion from this exercise is that child care represents the main component (or at least, a large component) of the expenditures on children in the PSID. In that sense, using time in non-parental child care as an input, and including expenditures on child care in the budget constraint of parents, can be viewed as focusing on the main component of expenditures on children and being specific about how it contributes to child skill accumulation. More specifically, I have money spent on the child in the form of child care affect child skill accumulation through changing how the child uses her time.

Table 24: Definition 1

	mean	sd	count
Ages [0,3]	0.67	0.29	84
Ages [0,5]	0.68	0.28	146
Ages [0,7]	0.71	0.26	223
Ages [0,9]	0.70	0.27	260
Ages [0,11]	0.70	0.27	275

Notes: Table 24 presents averages by age group for the fraction of total expenditure on children spent on child care. Definition 1 of total expenditures on children includes child care, money spent on toys, and money spent on school supplies

Table 25: Definition 2

	mean	sd	count
Ages [0,3]	0.55	0.28	84
Ages [0,5]	0.58	0.28	146
Ages [0,7]	0.61	0.26	223
Ages [0,9]	0.60	0.27	260
Ages [0,11]	0.60	0.27	275

Notes: Table 25 presents averages by age group for the fraction of total expenditure on children spent on child care. Definition 2 of total expenditures on children includes child care, money spent on toys, and money spent on school supplies.

Table 26: Definition 3

	mean	sd	count
Ages [0,3]	0.49	0.27	83
Ages [0,5]	0.52	0.27	144
Ages [0,7]	0.55	0.26	220
Ages [0,9]	0.53	0.26	256
Ages [0,11]	0.53	0.26	271

Notes: Table 26 presents averages by age group for the fraction of total expenditure on children spent on child care. Definition 3 of total expenditures on children includes child care, money spent on toys, money spent on school supplies, money spent on vacations, and money spent on clothes.

Table 27: Definition 4

	mean	sd	count
Ages [0,3]	0.42	0.24	71
Ages [0,5]	0.43	0.24	126
Ages [0,7]	0.46	0.24	194
Ages [0,9]	0.44	0.23	229
Ages [0,11]	0.43	0.23	243

Notes: Table 27 presents averages by age group for the fraction of total expenditure on children spent on child care. Definition 3 of total expenditures on children includes child care, money spent on toys, money spent on school supplies, money spent on vacations, money spent on clothes, and money spent on food.

G Assumptions About Time Inputs From Single Fathers

Empirically, single fathers usually don't raise children under five alone unless the mother has died; instead, they play a visiting role. In the ECLS-B there is a separate questionnaire for non-resident fathers, which is completed if the resident parent allows the survey-giver to contact the non-resident father, and the father agrees to participate. Besides concerns about egregious measurement error in the single father sample, as well as severe selection bias, there is also the problem of different variables being reported for the single father survey than on the resident parent questionnaire (in regards to activities done with children). Even if the single father sample were representative and the reporting accurate given the questionnaire single fathers receive, the available variables only allow an approximation of what is reported for resident parents.

The ECLS-B provides a non-resident father questionnaire (NRQ) in the first two waves of the survey. In this section, I document six points about the sample of non-resident fathers that complete this survey as well as attributes of single mothers in the data. Sample counts for these tabulations reflect response rates for the questions of interest; here, I am not restricting by whether I also observe variables necessary for the skill accumulation technology estimation. In the statistics presented below, I use survey weights for the primary caregiver sample in wave 2. The main purpose of this section is to establish that relatively few single fathers complete the survey, that those who do are not representative of the sample of single fathers, and that when they do complete the survey their answers and the answer's of their child's mother do not coincide (where comparable). In addition, Table 29 makes an additional point about the marital status composition of single mothers: most were never married. This coincides with the timing and nature of the marriage market in my model.

Selection in the Single Father Sample, Nature of Relationship with Child's Mother by Response Status

of SF The first three points are made in Tables 28 and 29. First, Table 28 shows that the response rate of non-resident fathers in each wave is about 1 in 3. Second, Table 29 shows that the marital status of the corresponding single mother is about the same for the group of families with a completed NRQ and without a completed NRQ. Third, Table 29 also shows that most single mothers were never married (about 70% and 65% in the first and second wave of the survey, respectively). Since I do not model divorce, the composition of marital status in single mothers is important to check.

Table 28: Response Rate NR Questionnaire

	(1) Wave 1	(2) Wave 2
Yes	0.300	0.309
No: Refusal	0.292	0.179
No: Not Permission	0.194	0.270
No: Ineligible, Lack of Contact	0.184	0.179
No: no NR	0.0290	0.0596
No: P not Biomother	0.000628	0.00303
Total	1	1
Obs.	2000	2000

Table 28 displays response rates of non-resident fathers to the non-resident father survey in the ECLS-B. Slightly less than one-third of non-resident fathers respond. Sample sizes rounded to nearest 50, following NCES requirements.

Table 29: Marital Status Composition of Mothers with NR fathers, by Questionnaire Response status

	Wave 1		Wave 2	
	(1) Completed NRQ	(2) No NRQ	(3) Completed NRQ	(4) No NRQ
Not Reported	0	0.00369	0	0.000118
Married	0.0640	0.0520	0.0921	0.0980
Separated	0.107	0.119	0.105	0.0939
Divorced	0.0909	0.0929	0.111	0.140
Widowed	0.00320	0.0142	0.00195	0.0169
Never Married	0.734	0.717	0.691	0.647
Not Biomother or Adoptive Parent	0	0.000897	0	0.00439
Total	1	1	1	1
Obs.	650	1350	650	1400

Table 29 displays the marital status composition of families where the biological parents are not cohabiting (single-parent families). The compositions are broken down by response status for the non-resident father questionnaire. Sample sizes rounded to nearest 50, following NCES requirements.

Single Father Visitation Frequency by NFQ Response Status, Influence of Single Father by NFQ Response Status SF vs. SM opinion The next three points are made in Tables 30-34. For point four, Table

30 tabulates the days since the non-resident father last saw the child. Fathers who complete the NRQ have seen the child on average 1.5 days more recently than fathers who do not. Fifth, in Table 30 I tabulate responses to the question “In a typical week, does [the child’s] father spend a lot, some, very little, or no time taking care of [the child]?”, for families without a completed NRQ’s (first column) and for those with an NRQ (second column). Fathers who completed the NRQ are almost 3 times more likely to be parenting with a resident primary caregiver who responds “A lot” to this question (35% compared to 12%). Relatedly, Table 32 shows that fathers who complete the NRQ are almost twice as likely to have seen their child in the last month than fathers who did not complete the NRQ (90% versus 46%). Sixth, in Tables 33 and 34 I tabulate the wave 2 responses to the question “When it comes to making major decisions, please tell me if [child’s] father has no influence, some influence, or a great deal of influence on such matters as child care?”, separately for mothers (Table 33) in families without an NRQ (column 1) and those with an NRQ (column 2) and fathers (Table 34) who completed the NRQ. Fathers who completed the NRQ think they have a lot of influence; mothers with children whose fathers completed the NRQ say they have less influence than the fathers claim, although they report more influence more than do mothers in families without a completed NRQ.

Table 30: Wave 1: Number of Days since NRF last saw child

	(1) No NRQ	(2) Completed NRQ
No. Days	3.860	2.353
Obs.	1300	650

Table 30 displays the average number of days since a non-resident father saw his child in the first wave of the survey, by response status to the non-resident father questionnaire. Sample sizes rounded to nearest 50, following NCES requirements.

Table 31: Wave 2: Frequency NRF last provides child care

	(1) No Completed NRQ	(2) Completed NRQ
Not Applicable	0.541	0.104
A lot	0.121	0.350
Some	0.117	0.280
Very little	0.0912	0.146
No time	0.130	0.120
Total	1	1
Obs.	1350	650

Table 31 displays the response to the question: “In a typical week, does [the child’s] father spend a lot, some, very little, or no time taking care of [the child]?”, for families without a completed NRQ’s (first column) and for those with an NRQ (second column). Sample sizes rounded to nearest 50, following NCES requirements.

Table 32: Wave 2: Number of Days since NRF last saw child

	(1) No NRQ	(2) Completed NRQ
Don’t Know	0.01	0
Refused	0.01	0
Not Applicable	0.06	0
Less than 1 month	0.459	0.896
More than 1 month, less than 1 yr	0.238	0.0718
More than 1 yr	0.0975	0.0198
No contact since birth/separation	0.133	0.0118
Total	1	1
Obs.	1350	650

Table 32 compares the amount of time since non-resident fathers last saw their child, by response status to the non-resident father questionnaire. Sample sizes rounded to nearest 50, following NCES requirements.

Table 33: Wave 2: Mother's Opinion of Father's Influence on CC

	(1) No NRQ	(2) Completed NRQ
Not Applicable	0.373	0.0207
No Influence	0.341	0.395
Some Influence	0.138	0.260
A Great Deal of Influence	0.148	0.324
Total	1	1
Obs.	1400	650

Table 33 tabulates mother's responses to the question: "When it comes to making major decisions, please tell me if [child's] father has no influence, some influence, or a great deal of influence on such matters as child care?", by response status for the non-resident father questionnaire. Sample sizes rounded to nearest 50, following NCES requirements.

Table 34: Wave 2: NRQ Father's Opinion of Father's Influence on CC

	Frequency
Not Ascertained	0.0318
No Influence	0.146
Some Influence	0.328
A Great Deal of Influence	0.494
Total	1
Obs.	650

Table 34 tabulates the response of father's who completed the non-resident father questionnaire to the question "When it comes to making major decisions, please tell me if you have has no influence, some influence, or a great deal of influence on such matters as child care?". Sample sizes rounded to nearest 50, following NCES requirements.